Administrative Records Mask Racially Biased Policing

DEAN KNOX   *Princeton University*
WILL LOWE   *Princeton University*
JONATHAN MUMMOLO   *Princeton University*

*R*esearchers often lack the necessary data to credibly estimate racial discrimination in policing. In particular, police administrative records lack information on civilians police observe but do not investigate. In this paper, we show that if police racially discriminate when choosing whom to investigate, analyses using administrative records to estimate racial discrimination in police behavior are statistically biased, and many quantities of interest are unidentified—even among investigated individuals—absent strong and untestable assumptions. Using principal stratification in a causal mediation framework, we derive the exact form of the statistical bias that results from traditional estimation. We develop a bias-correction procedure and nonparametric sharp bounds for race effects, replicate published findings, and show the traditional estimator can severely underestimate levels of racially biased policing or mask discrimination entirely. We conclude by outlining a general and feasible design for future studies that is robust to this inferential snare.

Word Count: 10285

Concern over racial bias in policing, combined with the public availability of large administrative data sets documenting police-civilian interactions, have prompted a raft of studies attempting to quantify the effect of civilian race on law enforcement behavior. These studies consider a range of outcomes including ticketing, stop duration, searches and the use of

This is a manuscript submitted for review.

force (e.g. Antonovics and Knight 2009; Fryer 2019; Nix et al. 2017; Ridgeway 2006). Most research in this area attempts to adjust for omitted variables that may correlate with suspect race and the outcome of interest. In contrast, this paper addresses a more fundamental problem that remains even if the vexing issue of omitted variable bias is solved: the inevitable statistical bias that results from studying racial discrimination using records that are themselves the product of racial discrimination (Rosenbaum 1984; Angrist and Pischke 2008; Elwert and Winship 2014). We show that when there is any racial discrimination in the decision to detain civilians—a decision which determines which encounters appear in police administrative data at all—then estimates of the effect of civilian race on subsequent police behavior are biased absent additional data and/or strong and untestable assumptions.

This paper makes several contributions. We clarify the causal estimands of interest in the study of racially discriminatory policing—quantities that many studies appear to be targeting, but are rarely made explicit—and show that the conventional approach fails to recover any known causal quantity in reasonable settings. Next, we highlight implicit and highly implausible assumptions in prior work and derive the statistical bias when they are violated. We proceed to develop informative nonparametric sharp bounds for the range of possible race effects, apply these in a reanalysis and extension of a prominent paper on police use of force (Fryer 2019), and present bias-corrected results that suggest this and similar studies drastically underestimate the level of racial bias in police-civilian interactions. Finally, we outline strategies for future data collection and research design that can mitigate these threats to inference. These are discussed in the context of a detailed and feasible proposed study of racial bias in traffic stops.

As we show in this paper, the difficulty of estimating racial bias using police records stems from a thorny combination of mediation (Hernán et al. 2004; Imai et al. 2011; Robins et al. 2000; Pearl 2001; VanderWeele 2009) and selection (Heckman 1979; Lee 2009): the effect of civilian race on the outcome of a police encounter is mediated by whether the civilian is stopped by police, but analysts only have data for one level of the mediator—i.e., data on stopped individuals. Because of this, police records do not contain a representative sample of all individuals that police observe, but rather only those civilian encounters which escalated to the point of triggering a reporting requirement. If a civilian's race affects whether officers choose to stop that civilian (Gelman et al. 2007; Glaser 2014),

then analyzing administrative police records amounts to conditioning on a variable that is itself affected by suspect race, namely, whether a suspect appears in the data at all. This could occur if officers have a higher threshold for stopping white civilians during the unseen first stage of police-civilian contact, meaning that white civilians observed in the data are incomparable because they tend to pose a greater threat to police than observed minorities. These unobserved differences can lead analysts to understate anti-minority racial bias—or even produce the appearance of anti-white bias—in the use of force. Despite claims to the contrary (Fryer 2018, 2), this statistical bias often cannot be eliminated with additional control variables, even if the goal is to estimate causal effects among the subset of police-civilian encounters that appear in police data. Moreover, the problem remains whether racial bias in detainment stems from so-called "taste-based" or "statistical" discrimination (Arrow 1972, see below for extended discussion on this point).

At first glance, the problem of race-based selection into policing data may appear a classic case of sample selection bias (Elwert and Winship 2014; Heckman 1979) for which numerous remedies already exist. But policing data exhibit a constellation of features that render previous methodological approaches unsuitable or unusable in this setting, leading prominent scholars in this area to declare that "It is unclear how to estimate the extent of such bias or how to address it statistically," (Fryer 2018, 5).[1] For example, path-breaking work including Heckman (1979) and more recent extensions like Lee (2009), provide methods for estimating or bounding average treatment effects in the population while accounting for sample selection. But with only data on stopped individuals, policing scholars rarely seek to estimate population treatment effects, instead targeting effects among individuals who actually interact with police. We show that even without attempting to generalize to the broader population, the issues we raise result in biased estimates of the effect of race on police behavior *even among encounters in which civilians are detained*.

A related large literature provides remedies for so-called "post-treatment bias"—statistical bias which results from conditioning on a variable that is affected by the causal variable of interest

---

[1]This comment was made in reference to an analysis of arrest data in (Fryer 2019). Further, (Fryer 2019) includes an analysis aimed at characterizing selection into police data sets, and finds mixed results depending on the outcome examined. The study states: "Taken together, this evidence demonstrates how difficult it is to understand whether there is potential selection into police data sets . . . Solving this is outside the scope of this paper," (19).

(Rosenbaum 1984). But implementation of these techniques requires either knowledge of the scale of the missing data (e.g. Nyhan et al. 2017), or complete data on the post-treatment variable (e.g. Acharya et al. 2016).[2] In the case of policing, administrative data sets only include observations with one level of the post-treatment variable (i.e. data on stopped individuals) *and* give no purchase on the number of individuals police observe but do not stop, meaning these techniques cannot be applied. This scenario also differs from situations of "truncation by death" (Frangakis and Rubin 2002) in which receipt of a treatment causes sample attrition and renders outcomes for some portion of units undefined. In the policing setting, individuals not detained by police are absent from the data, but many outcomes of interest are often still defined (e.g. the level of force applied to non-stopped individuals is zero, a realized outcome). This feature allows us to identify additional causal quantities that cannot be recovered in the "truncation by death" setting. In short, existing methods offer either unusable or sub-optimal solutions to this pernicious threat to inference, absent strong assumptions about the unseen process mapping civilian race to officers' decisions to detain individuals.

Our analysis indicates that existing empirical work in this area is producing a misleading portrait of evidence as to the severity of racial bias in police behavior. Replicating and extending the study of police behavior in New York in Fryer (2019), we show that the consequences of ignoring the selective process that generates police data are severe, leading analysts to dramatically underestimate or conceal entirely the differential police violence faced by civilians of color. For example, while a naïve analysis that assumes no race-based selection into the data suggests only 10,000 black and Hispanic civilians were handcuffed due to racial bias in New York City between 2003 and 2013, we estimate that the true number is approximately 56,000. And while analyses ignoring bias in stopping would conclude that 10% of uses of force against black and Hispanic civilians in these data were discriminatory, after bias-correction, we estimate that the true percentage is 39%.

While the techniques used to obtain our corrected results eliminate several facially implausible (and in some cases, empirically falsified) assumptions that are implicit in prior work, we caution that they nevertheless rely on weaker assumptions that in some cases are difficult to verify, as we discuss

---

[2]In addition, the remedy proposed in Blackwell (2013), which requires re-weighting across all strata of the post-treatment variable, cannot be implemented in the situation we describe. However, the alternative designs we propose below are amenable to this approach.

below. We seek to advance the study of racial bias in policing by explicitly stating these assumptions, discussing their plausibility in this context, and carefully grounding unobservable parameters—in particular, the proportion of racially discriminatory minority stops, which relates closely to the severity of the statistical bias—in prior research (Goel et al. 2016; Gelman et al. 2007). We show that obtaining more precise bias-corrected estimates of racial discrimination in policing requires future research to be designed with this issue in mind. To that end, we outline a research design that alleviates these concerns. Our study also provides a general framework for analyzing the study of racial bias that can illuminate the causal interpretation of other longstanding tests for discrimination. For example, we show that under reasonable assumptions, so-called "outcome tests," which compare the rates of finding evidence of criminal activity across detained suspects of different racial groups (Knowles et al. 2001), imply a lower bound on the share of racial minorities who are discriminatorily detained. Outcome tests also appear elsewhere in criminal justice studies, e.g. in capital sentencing (Alesina and Ferrara 2014) and bail decisions (Arnold et al. 2018). And as Ayres (2002) and Simoiu et al. (2017) note, such tests have also been applied in a range of other social contexts, including financial lending and editorial decisions. By nesting the study of discrimination in a rigorous and general causal framework, our study can help synthesize results from a broad interdisciplinary literature on racial bias.

Our work also extends a growing literature in political science examining the political implications of law enforcement which, in recent decades, has largely studied policing indirectly, e.g. as a means of explaining political participation (Burch 2013; Cohen et al. 2017; Lerman and Weaver 2014; White 2019) or as an instance of bureaucracy (Brehm and Gates 1999; Lipsky 1980; Ostrom and Whitaker 1973; Wilson 1989). This work is pathbreaking, but with some recent exceptions (Harvey and Mungan 2019; Magaloni et al. 2015; Mummolo 2018a; Peyton et al. 2019; Soss and Weaver 2017), has tended to conceptualize policing as a cause of politics, rather than a political act in and of itself. The field's relative inattention to policing was made evident by several recent officer-involved shootings of unarmed black men (Edwards et al. 2019) and subsequent social unrest that caught many political scientists flatfooted, with little systematic evidence to offer as the demand for explanations of police behavior surged. As Soss and Weaver (2017) notes, the field's limited store of relevant knowledge in the aftermath of these events was especially glaring given law enforcement's role as an everyday

conduit of state power. According to one often-cited definition, politics is "who gets what, when, how" (Lasswell 1936). As a matter of routine, the dynamics of police-civilian interactions determine who gets protected, punished, or left to fend for themselves (Wilson 1968). Viewed in this way, the role of race in the state's exercise of violence, and in the provision of safety more broadly, is inherently political (Alexander 2010; Gottschalk 2008; Key 1949). In addition to offering a rigorous analytic framework to help researchers contend with longstanding methodological hurdles, our study also underscores an often overlooked truth: policing is high-stakes politics.

CONCEPTUALIZING RACE AS A CAUSAL VARIABLE

We regard the investigation of racial bias in policing as an inherently causal inquiry, albeit a notoriously difficult one. That is, researchers seek to assess whether police behavior during police-civilian encounters would have differed if the civilian had belonged to another racial group, holding constant civilian behavior and circumstance. As noted in Fryer (2018), this "'race effect'... is the proverbial 'holy grail' — the parameter that we are all attempting to estimate but never quite do" (2). This task is distinct from the descriptive enterprise of merely documenting differential police behavior during encounters with various groups, as such disparities can arise via numerous processes that do not imply racial discrimination.[3]

The notion of a "causal effect of race" on an individual's outcome is the subject of much contention in the literature on causal inference (Hernán 2016; Pearl 2018). Most notably, some have argued that this effect is undefined because race is an immutable, and hence non-manipulable, characteristic (Holland 1986). Others argue that an individual's race is a complex, multifaceted treatment—a "bundle of sticks," in the words of Sen and Wasow (2016)—that affects outcomes through myriad channels, and therefore researchers must be precise about the specific facets of race under consideration (Greiner and Rubin 2011).

---

[3]For example, we may observe that members of one racial group are stopped more often by police than members of another racial group. While this result shows disparate police behavior, it does not conclusively demonstrate that the difference is due to civilian race. It could simply be the case that members of the first group participate in criminal activity more often in public.

Our analysis avoids this debate by focusing on police-civilian *encounters*—i.e., sightings of civilians by police—as the unit of analysis, rather than individuals. The manipulation of race is conceptualized as the counterfactual substitution of an individual with a different racial identity into the encounter, while holding the encounter's objective context—location, time of day, criminal activity, etc.—fixed. In other words, the "treatment" in this case is the entire "bundle of sticks" encapsulating the race of the civilian—including skin tone, dialect, and clothing. We note that the credibility of causal inferences and the exact interpretation of racial discrimination in this framework will depend crucially on how the analyst defines "race." We leave the specific operationalization in a given context to the analyst, and, in line with advice in Sen and Wasow (2016), encourage scholars to carefully convey their conceptualization of race when studying this and related questions.[4]

By conceptualizing the treatment in this way, we avoid consideration of the perhaps implausible counterfactual of holding all features *of an individual* constant but for their race. While various aspects of racial identity and its close correlates may not be separable in the observed world, there exists a subset of comparable *situations* in which minority and majority civilians are observed by police. If this subset can be identified, or approximated through covariate adjustment, we can estimate the counterfactual police behavior that would have occurred had the civilian in question been replaced with a member of another racial group.

While our approach considers a valid counterfactual and isolates racial discrimination that occurs during police-civilian encounters, it necessarily mutes the influence of pre-encounter macro-institutional factors, such as decisions to deploy more officers to communities of color. In keeping with the goals of prior studies in this area, our approach holds such contextual features constant, allowing us to ask whether an encounter would have unfolded differently had it involved a civilian of differing race. But even if no such difference exists within encounters, law enforcement strategies adopted before encounters occur could still produce racially biased policing. We caution readers to keep this scope condition in mind.

---

[4]Note that while the unit of analysis is the police-civilian encounter, for the sake of brevity, we occasionally refer to "minority civilians" as shorthand for "police-civilian encounters with minority civilians" in subsequent discussion. Readers are cautioned to keep this distinction in mind.

PRIOR RESEARCH ON RACIAL BIAS IN POLICING

Race-based selection into policing data has been previously noted, and some scholars have devised research designs in an attempt to sidestep this issue. Grogger and Ridgeway (2006), for example, leverage the so-called "veil of darkness" strategy, comparing patterns in traffic stops that occur before and after sunset under the logic that the race of the driver is plausibly hidden to police officers after dark. In this way, the study aims to identify a sample of police-civilian interactions that were initiated in a race-blind manner. Similarly West (2018) examines data on police responses to traffic incidents, arguing that whether a co-racial officer responds to a motorist's unanticipated accident is as-if random. If the assumptions in these studies hold, concerns over race-based sample selection are greatly alleviated.

These attempts to mitigate race-based selection remain rare, as most empirical studies in this literature focus nearly exclusively on mitigating the more familiar problem of omitted variable bias. For example, Fryer (2019) (detailed below), a study of racial bias in police violence, estimates discrimination using data on police-civilian encounters via multivariate regressions that control for a host of observables relating to civilians, officers and circumstance. In a related article, the author asserts that "regression can recover the 'race effect' if race is 'as good as randomly assigned,' conditional on the covariates" (Fryer 2018, 2). Fryer (2019) claims to find evidence of bias in sub-lethal force but none in lethal encounters.

A related study, Johnson et al. (2019), attempts to estimate racial bias in police shootings. Examining only positive cases in which fatal shootings occurred, they find that the majority are white and conclude from this that no anti-minority bias exists. Knox and Mummolo (ming) show that this conclusion rests on the erroneous assumption that police encounter minority and white civilians in equal number.

Prior work has also examined racial bias during traffic stops, such as Ridgeway (2006) which employs propensity score weighting when estimating racial bias in traffic stops in Oakland, CA. The analysis examines outcomes including citations, stop duration and the decision to search cars. The study claims this re-weighting strategy can recover "the causal effect of race" (9) on post-stop outcomes. In general the analysis finds little evidence of racial bias on most outcomes, with the exception of stop duration. Antonovics and Knight (2009) uses data on traffic citations from the Boston Police Department to estimate the probability that a ticketed driver was searched, controlling for driver

attributes such as age, race and gender as well as neighborhood traits. They interpret the coefficient on an indicator of whether the officer and ticketed driver are of different races as an estimate of "racial profiling based upon prejudice," as opposed to statistical discrimination (167). The claim is implicitly causal: some share of searches among racially mismatched driver-officer pairs would not have occurred had the driver belonged to another racial group.

The above examples represent a mere fraction of a decades-long, multi-disciplinary effort to quantify the degree to which police discriminate against civilians of color (see Fridell (2017), Atiba Goff and Kahn (2012), and Ridgeway and MacDonald (2010) for more extensive reviews of this empirical literature). We highlight these specific examples because they all contain several common features that are central to our critique. For one, all of these studies analyze data that fail to capture the unseen selective process through which police come to engage civilians, a process that prior work shows is function of civilian race (Gelman et al. 2007). In this way, these studies all fail to account for the impact of race on the composition of the sample under study. As we show below, failing to account for this undocumented first stage of the police-civilian interaction will lead to statistical bias, even if the goal is to estimate the effect of suspect race within the sample of individuals who appear in police data and, in many cases, even with a "complete" set of control variables that render civilian race as-if randomly assigned to police encounters.

Second, all of the aforementioned studies, despite making at least implicitly causal claims, leave ambiguous the precise quantity of interest—whether it be the average treatment effect (ATE) of race in all encounters; the average treatment effect among the subset of encounters appearing in police data because a stop was made ($\text{ATE}_{M=1}$), which differs tremendously from the ATE; or the markedly more restrictive and difficult-to-interpret controlled direct effect among the same subset ($\text{CDE}_{M=1}$, defined below). While studies commonly discuss omitted variable bias and attendant assumptions, they rarely discuss the additional assumptions necessary to identify specific causal quantities of interest. As a result, readers are unable to assess the adequacy of research designs and estimators, rendering the interpretation and policy relevance of much prior work ambiguous.

Taste-based vs. statistical discrimination

A closely related literature attempts to parse "taste-based discrimination" (racial animus) from so-called "statistical discrimination" (Arrow 1972, 1998; Becker 1971; Eberhardt et al. 2004; Phelps 1972) as mechanisms for racially biased policing, and instead focus on recovering the causal effect of civilian race on police behavior. In this paper, we do not attempt to disentangle these mechanisms, and we note that taste-based and statistical discrimination both pose serious normative concerns. While statistical discrimination is sometimes viewed as more innocuous, it nonetheless constitutes racial profiling since officers detain civilians not due to their own observed actions but rather those of the racial groups to which civilians belong. Federal courts have outlawed such justifications for detaining civilians (see discussions in Goel et al. (2016) and Mummolo (2018b)). Thus, quantifying the causal effect of civilian race on police behavior—our task here—is imperative regardless of the mechanism that produces such an effect.

CLARIFYING THE EFFECT OF CIVILIAN RACE: NOTATION, ESTIMANDS, ASSUMPTIONS, AND EXISTING APPROACHES

Researchers and policymakers examining the effects of racially biased policing are nominally interested in the relationship between two variables: the race of the civilian involved in encounter i, which we operationalize through their minority status $D_i \in \{0, 1\}$, and consequent police behavior Y_i. However, analyses of administrative data on police-civilian encounters inherently involve a mediating variable that may be affected by race: whether an individual is stopped by police, which we denote M_i. The causal ordering of these variables is depicted in the directed acyclic graph (DAG) in Figure 1. We note that analysts often possess rich contextual information about the objective context of the encounter, such as its location and time, which may relate to all of the above. We denote these covariates collectively as X_i. However, administrative data invariably fail to capture unobservable subjective aspects of the encounter, U_i, such as an officer's suspicion or sense of threat.

As a motivating example, we consider the challenge of estimating racial bias in police violence as recently attempted in Fryer (2019). We ground our analysis in the potential outcomes framework

FIGURE 1. **Directed acyclic graph of racial discrimination in the use of force by police. Observed X is left implicit; these covariates may be causally prior to any subset of D, M, and Y.**

(Rubin 1974) often used in the study of causal mediation (Pearl 2001; Imai et al. 2011). The potential mediator $M_i(d)$ represents whether encounter i would have resulted in a stop if the civilian were of race d. Similarly, the potential outcome $Y_i(d, m)$ represents whether force would have been used in encounter i if the civilian were of race d and the mediating variable were m. The observed mediator and outcome can be written in terms of these potential values as $M_i = M_i(D_i) = \sum_d M_i(d)\mathbf{1}\{D_i = d\}$ and $Y_i = Y_i(D_i, M_i(D_i)) = \sum_d \sum_m Y_i(d, m)\mathbf{1}\{D_i = d, M_i = m\}$, respectively. For any individual encounter, the (unobservable) causal effect of civilian race is the difference in potential force if the civilian were a minority and stopped as if they were a minority, versus if they were white and stopped accordingly, $Y_i(1, M_i(1)) - Y_i(0, M_i(0))$.

This notation implicitly makes the stable unit treatment value assumption (SUTVA, Rubin 1990). "Stability" is of particular note: this stipulates that finer racial gradations must not affect the way that officers behave, *above and beyond* any differences between the broad binary categories $D_i = 0$ and $D_i = 1$. SUTVA also requires that each encounter is unaffected by a civilian's race in other encounters; this might be violated if, for example, groups of individuals are stopped simultaneously.

Traditionally, analysts use data on stopped individuals to study bias by computing the difference in violence rates between stopped minority and white civilians, while controlling for observable differences between these two sets of encounters. We term this the "naïve estimator," $\hat{\Delta}$, and it can be written as:

$$\hat{\Delta} = \overline{Y_i | D_i = 1, M_i = 1} - \overline{Y_i | D_i = 0, M_i = 1} \tag{1}$$

where conditioning on possible treatment-outcome confounders, X_i, is left implicit. Assuming the analyst has correctly measured and specified all such confounders, $\hat{\Delta}$ may appear entirely reasonable at

first glance. However, without further assumptions, this quantity will have no causal interpretation so long as the treatment affects the mediator (i.e. civilian race affects whether officers detain a civilian). As we show below, this is because encounters with treated (minority) civilians that result in a stop ($M_i = 1$) will not be comparable to those with stopped control (majority) civilians. As a simple example, suppose officers exhibited racial bias as follows: they detain white civilians if they observe them committing a serious crime (like assault, potentially warranting the use of force) but detain nonwhite civilians regardless of observed behavior. When this is true, comparing stopped white and nonwhite civilians amounts to comparing fundamentally different groups. The analyst will observe force used against a greater proportion of stopped white civilians because of the unobserved severity of their crimes.[5] Under the traditional approach, the analyst would naïvely conclude that anti-white bias exists, yielding an erroneous portrait of racial discrimination in the use of force.

To formalize the limitations of the naïve estimator, we begin by partitioning the population into principal strata with respect to the mediator (Frangakis and Rubin 2002; VanderWeele 2011). That is, we conceptualize police-civilian encounters in terms of four latent classes within which $M_i(1)$ and $M_i(0)$ are constant. The general approach of principal stratification has proven useful for clarifying and bounding quantities of interest in areas ranging from instrumental variables (Angrist et al. 1996; Balke and Pearl 1997) to the closely related "truncation by death" problem (Rubin 2000; Zhang and Rubin 2003).

These principal strata include "always-stop" encounters in which $M_i(0) = M_i(1) = 1$, as well as racially discriminatory stops ("racial stops") in which $M_i(1) = 1$ but $M_i(0) = 0$. Always-stop encounters may be conceptualized as relatively severe scenarios, such as violent crimes in progress, in which officers have no choice but to intervene regardless of civilian race. In contrast, previous work has identified certain behaviors, such as "furtive movements" (Gelman et al. 2007; Goel et al. 2016), that appear to be acted upon selectively by officers based on the race of suspects. "Never-stop" encounters, where $M_i(0) = M_i(1) = 0$, are situations in which civilians appear inconspicuous and would not be stopped, regardless of race. There also may be anti-white racial encounters, in which

---

[5]While some police records indicate whether a suspect was engaged in violent behavior, allowing the analyst to control for this particular factor, a host of similar concerns (e.g. time-varying officer suspicion) are unmeasured and thus cannot be controlled away.

12

$M_i(1) = 0$ but $M_i(0) = 1$, though we believe these to be rare to nonexistent (discussed further below). Figure 2 shows encounters appearing in police records (principal strata for which $M_i(D_i) = 1$) are not comparable across civilian races. Minority police-civilian encounters that result in a stop are a mixture of "always-stop" and "anti-minority racial stop" encounters, while encounters with white civilians that result in a stop are a combination of "always-stop" and "anti-white racial stop" encounters. These are fundamentally different groups, and without further assumptions, comparisons of rates of violence between them using the naïve estimator will be statistically biased.
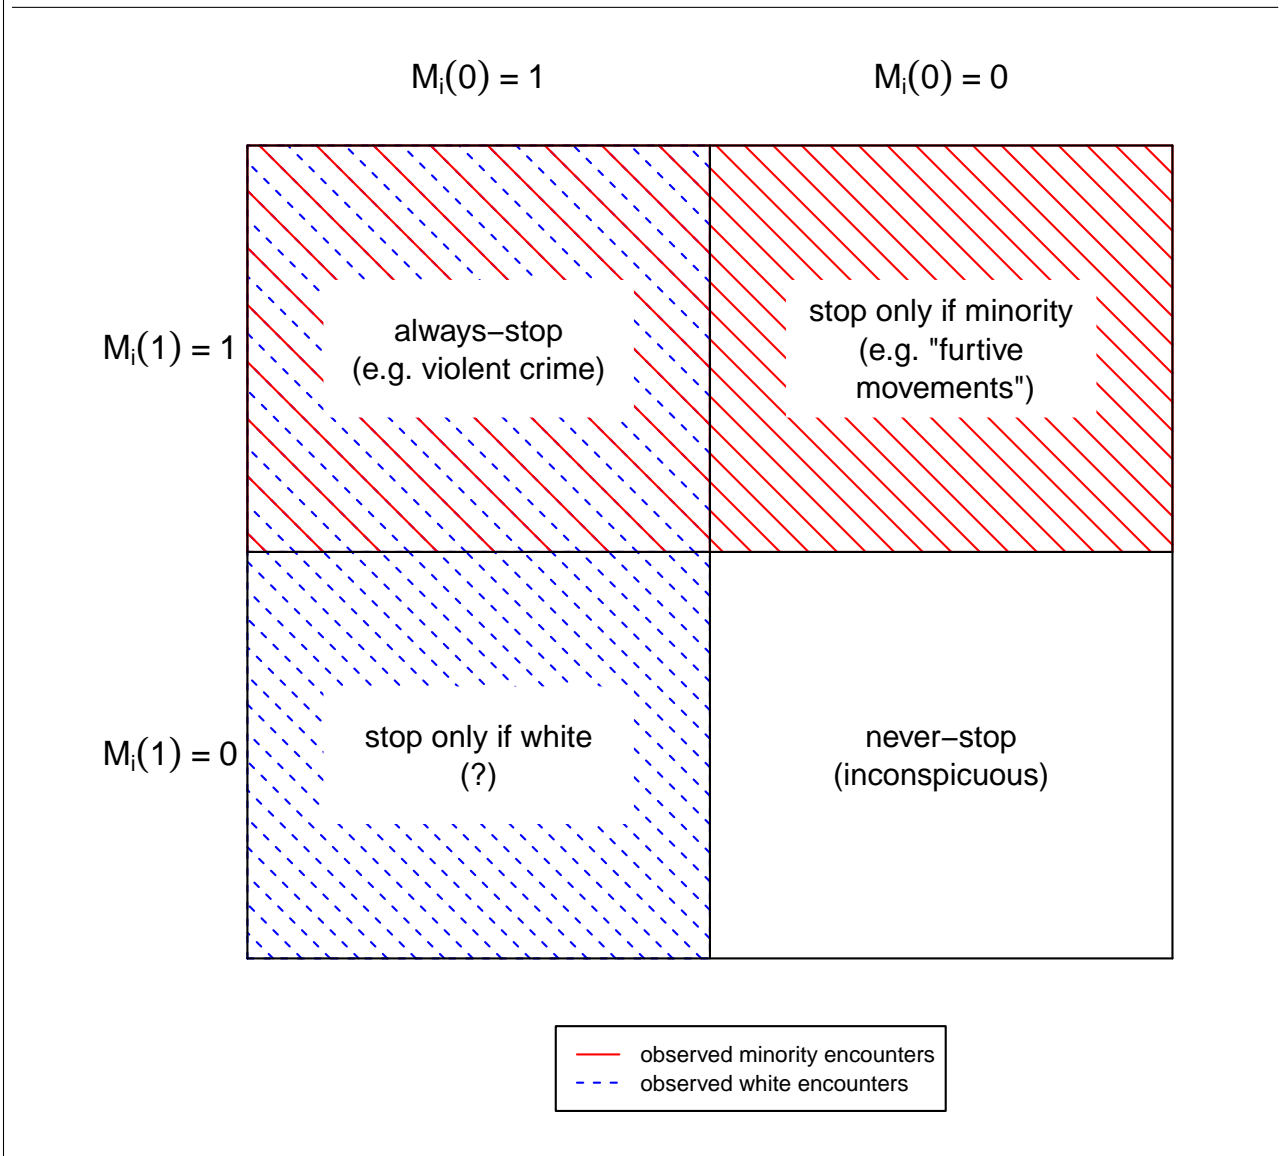
To state this more formally, note that the naïve estimator recovers the weighted combination of violence rates in observed principal strata:

$$\mathbb{E}[\hat{\Delta}] = \mathbb{E}[Y_i | D_i = 1, M_i = 1] - \mathbb{E}[Y_i | D_i = 0, M_i = 1]$$

$$= \mathbb{E}[Y_i(1,1) | D_i = 1, M_i(1) = 1, M_i(0) = 1] \Pr(M_i(0) = 1 | D_i = 1, M_i(1) = 1)$$

$$+ \mathbb{E}[Y_i(1,1) | D_i = 1, M_i(1) = 1, M_i(0) = 0] \Pr(M_i(0) = 0 | D_i = 1, M_i(1) = 1)$$

$$- \mathbb{E}[Y_i(0,1) | D_i = 0, M_i(1) = 1, M_i(0) = 1] \Pr(M_i(1) = 1 | D_i = 0, M_i(0) = 1)$$

$$- \mathbb{E}[Y_i(0,1) | D_i = 0, M_i(1) = 0, M_i(0) = 1] \Pr(M_i(1) = 0 | D_i = 0, M_i(0) = 1). \tag{2}$$

In Equation 2, the first term is the average rate of force applied during encounters with racial minorities of the always-stop stratum, while the second term deals with minorities in the anti-minority racial-stop stratum. The third and fourth terms are the average violence rates among *white* civilian encounters in the always-stop and anti-white racial stop strata. Importantly, principal strata are not fully observable without further assumptions, and they exist even after conditioning on X_i: for any particular minority stop, it is fundamentally impossible to know with certainty whether a white civilian would have been stopped in identical circumstances. In sum, the naïve estimator compares groups with different potential outcomes, and because these groups are unobservable, the resulting bias is difficult to address.

A central quantity of interest in the study of policing bias is the average treatment effect of race, ATE $= \mathbb{E}[Y_i(1, M_i(1)) - Y_i(0, M_i(0))]$—the extent to which civilians of color face greater risk of police violence than white civilians *because of their race*. The ATE considers both reported and unreported encounters, and it captures two related phenomena: first, whether members of the minority are differentially stopped; and second, if they are differentially subject to violence. However, police

FIGURE 2. Principal Strata and Observed Police-Civilian Encounters. The figure displays the four principal strata that comprise police-civilian encounters based on how the mediator M (whether a civilian is stopped by police) responds to treatment D (whether the civilian is a racial minority). Minorities in the "always stop" and anti-minority racial stop strata, highlighted in red, are stopped by police, and thus appear in police administrative data. Likewise, white civilians in the "always stop" and anti-white racial stop strata, highlighted in blue, appear in police data. "Never stop" encounters are unobserved. Because white and nonwhite encounters are drawn from different principal strata, the two groups are incomparable and estimates of causal quantities using observed encounters will be statistically biased absent additional assumptions.

	$M_i(0) = 1$	$M_i(0) = 0$
$M_i(1) = 1$	always–stop (e.g. violent crime)	stop only if minority (e.g. "furtive movements")
$M_i(1) = 0$	stop only if white (?)	never–stop (inconspicuous)

——— observed minority encounters
- - - observed white encounters

administrative records contain data only on reported encounters, meaning that this quantity cannot be estimated solely with police administrative data without untenable assumptions. The ATE can be written as:

$$
\begin{aligned}
\text{ATE} &= \mathbb{E}[Y_i(1, M_i(1))] - \mathbb{E}[Y_i(0, M_i(0))] \\
&= \sum_d \sum_m \sum_{m'} \Big(\mathbb{E}[Y_i(1, M_i(1)) - Y_i(0, M_i(0))|D_i = d, M_i(1) = m, M_i(0) = m'] \\
&\qquad\qquad \times \Pr(D_i = d, M_i(1) = m, M_i(0) = m') \Big),
\end{aligned}
\tag{3}
$$

where the second line illustrates how it sums over the principal strata depicted in Figure 2, taking into account the number of minority and white civilians in each strata (the probabilities) and the local average treatment effects for each group (the expectations). In Online Appendices A.1–A.4, we use these quantities to derive bias and nonparametric sharp bounds.

No data is available for "never-stop" encounters, those with $M_i(1) = M_i(0) = 0$. Moreover, racial-stop encounters, with $M_i(1) = 1$ and $M_i(0) = 0$, are only recorded for minority civilians. However, consistent with (Nyhan et al. 2017), we show in Online Appendix A.6 that the ATE can be point identified if researchers collected two additional numbers: the count of total minority and white encounters, within levels of covariates X where applicable—a point we discuss further in our recommendations for future research.[6]

Because "never-stop" encounters are unobserved in current data sources, researchers seeking to understand the role of race in police behavior have, at least implicitly, focused on more narrowly defined

---

[6]Nyhan et al. (2017) examines the problem of studying the effect of party identification on turnout using voter registration files, given the fact that party ID likely affects who registers to vote. In an approach that is equivalent to our Proposition 2, the study uses registration rates of the treated and control voting-age populations to bound the ATE given the effect of party ID on registration. This option is not available in practice here since no data sets contain information on unreported encounter rates, or even their order of magnitude. As a result, analysts in this literature focus almost exclusively on the $\text{ATE}_{M=1}$, which we examine with a different approach here. While some work in policing uses population figures as proxies for these encounter rates, we are skeptical of this approach, as police frequently stop civilians who reside in other jurisdictions.

estimands.[7] Studies commonly restrict analysis to the subset of reported encounters, i.e., they seek to estimate effects among those stopped by police, $\text{ATE}_{M=1}$. In contrast with the ATE, this estimand is by definition not concerned with unreported white encounters that *would have* escalated to a stop if the involved civilian was a minority. (The same is true for unreported black encounters that *would have* escalated if the involved civilian was white, to the extent that this group exists.) Formally, this quantity is given by

$$\text{ATE}_{M=1} = \mathbb{E}[Y_i(1, M_i(1))|M_i = 1] - \mathbb{E}[Y_i(0, M_i(0))|M_i = 1]. \tag{4}$$

Relatedly, analysts may seek to causally attribute the number of minority stops in which force would not have been used if the individual in question had been white (Yamamoto 2012). This value is proportional to the conditional average treatment effect among the treated (i.e., minority) stops, which can be written as:

$$\text{ATT}_{M=1} = \mathbb{E}[Y_i(1, M_i(1))|D_i = 1, M_i = 1] - \mathbb{E}[Y_i(0, M_i(0))|D_i = 1, M_i = 1] \tag{5}$$

While the average treatment effects are of obvious policy importance, they are not the only quantity that researchers might seek to estimate. A closely related estimand is the controlled direct effect among the subset of reported encounters, $\text{CDE}_{M=1} = \mathbb{E}[Y_i(1, 1)|M_i = 1] - \mathbb{E}[Y_i(0, 1)|M_i = 1]$. This estimand differs from the $\text{ATE}_{M=1}$ in its conceptual approach to racially discriminatory stops. Where the $\text{ATE}_{M=1}$ asks whether a stop would have occurred *at all* if the individual were of differing race, the $\text{CDE}_{M=1}$ seeks to quantify what would have happened *if the officer was forced to stop them anyway*, perhaps against the officer's will. In practice, the difference is one of interpretation—regardless of the target quantity, existing work in this domain is based on the naïve difference in reported outcomes, and the question lies in the interpretation of estimated results. We note that causal estimands in the literature are often left undefined, making it difficult to assess whether published results are intended to correspond to the $\text{ATE}_{M=1}$ or $\text{CDE}_{M=1}$ (e.g. Goel et al. 2016; Simoiu et al. 2017). In Online Appendix A.3, we discuss

---

[7]For example, Fryer (2018) notes that his analysis of police use of force is estimating the effect of suspect race "conditional on an interaction," with police (4), rather than seeking its average treatment effect in the population.

the $CDE_{M=1}$ at length. We show that it cannot be recovered in this setting unless analysts make the untenable assumption that no mediator-outcome confounding exists (Assumption 5, below). We refer readers to the Online Appendix for further details, and focus on recovery of average treatment effects here.

Necessary assumptions

In this subsection, we describe a number of statistical assumptions that the analyst must make for a causal study of racially biased policing when only administrative data on police-civilian interactions is available. Without these assumptions, causal quantities of interest in this substantive area cannot be identified in data.

Assumption 1 (Mandatory Reporting). $Y_i(d, 0) = 0$ *for all i and for* $d \in \{0, 1\}$.

We assume all encounters that escalate to the use of force also trigger a reporting requirement and are therefore observed in administrative data. Though there exists wide variability in data recording practices across jurisdictions, this assumption is plausible in the study of many major police departments. For example, New York Police Department (NYPD) officers are required to report a number of variables, including the specific type of force used, following each "stop, question, and frisk" encounter. Based on these and other reports, the NYPD releases detailed annual use-of-force reports (NYPD 2017). The completeness of these reports with respect to fatalities is informally enforced by standard journalistic practices which place high emphasis on documenting incidents of violent crime (Iyengar 1994). Lesser forms of force are more likely to go unreported, to be sure, but the ubiquity of surveillance cameras, cell phone cameras, and media interest in police brutality makes unobserved uses of force increasingly unlikely (Fisher and Hermann 2015). We note that this assumption is implicit in all analyses of police use of force that rely on administrative data.

Assumption 2 (Mediator Monotonicity). $M_i(1) \geq M_i(0)$ *for all i*.

This assumption allows that there may be encounters in which minorities would be stopped ($M_i(1) = 1$) but whites would not ($M_i(0) = 0$), perhaps because officers racially discriminate in applying differential thresholds of "reasonable suspicion." However, we assume that the reverse is

never true: white civilians are never stopped in circumstances when their minority counterparts would be allowed to pass. This is clearly a stylized representation of a complex reality, and it would be violated if minority officers discriminate against white civilians. This could also occur if white civilians were more likely to be stopped by police because they appeared out of place in a predominantly black neighborhood, perhaps under the assumption that they were there to buy drugs (Gelman et al. 2007, 822). These are rare occurences, and as a robustness check to probe the validity of this assumption, Figure B4, our reanalysis of Fryer (2019) after dropping all stops based on suspicion of a drug transaction, shows substantively similar results.

Assumption 3 (Relative Non-severity of Racial Stops)**.**

$$\mathbb{E}[Y_i(d,m)|D_i = d', M_i(1) = 1, M_i(0) = 1, X_i = x] \geq \mathbb{E}[Y_i(d,m)|D_i = d', M_i(1) = 1, M_i(0) = 0, X_i = x]$$

We theorize that for encounters during criminal events severe enough to warrant stopping a civilian regardless of race (i.e., "severe" or "always-stop" encounters), the use of force is as or more likely to occur than during encounters in which police have more discretion over whether to stop an individual (i.e., those in which racial discrimination in stopping can occur) in expectation. We regard this assumption, which compares violence rates within encounters that hold civilian race fixed, as highly plausible. As one hypothetical example, this assumption would imply that police are as or more likely to use force against a white civilian observed committing assault than a white civilian observed jaywalking, on average.

Assumption 4 (Treatment Ignorability)**.**

(a) *With respect to potential mediator* $M_i(d) \perp\!\!\!\perp D_i | X_i$.

(b) *With respect to potential outcomes:* $Y_i(d,m) \perp\!\!\!\perp D_i | M_i(0) = m', M_i(1) = m'', X_i$.

This states that conditional on X_i, civilian race is "as good as randomly assigned" to encounters, and officers encounter minority civilians in circumstances that are objectively no different from white encounters. Part 4(a) stipulates that the observed covariates X include the confounder W in Figure 3(a). This assumption, while strong, has become more plausible in recent years as administrative data sets have come to include a host of encounter attributes that might largely capture features observable to police which correlate with suspect race and the potential for force. However, we note that this

FIGURE 3. Violations of assumptions. DAGs (a), (b), and (c) respectively illustrate the violation of Assumptions 4(a), 4(b), and 5. Note that the variable U depicted in DAG (c) is almost certain to exist in the policing context, and we do not advocate the use of Assumption 5.

cannot be tested, even indirectly, without data on non-stopped individuals. This assumption would be violated if neighborhoods with high shares of minority residents were more heavily policed and the analyst failed to adjust for neighborhood, e.g. using fixed effects. Part 4(b) implies that, for example, if police were more heavily armed during minority-neighborhood patrols and hence more likely to deploy force—represented by V in Figure 3(b)—then V must be included in X. Without Assumption 4, the range of possible racial effects is so wide as to be uninformative. We also note that every study claiming to estimate racial discrimination using similar data makes this assumption, often implicitly. Our aim in this paper is not to assert the plausibility of treatment ignorability, but rather to clarify that deep problems remain even if this well-known issue is somehow solved.

Strong assumptions

We now discuss further assumptions that are often left implicit in empirical studies of racially biased policing, and that are implausible in many settings. We illustrate these scenarios graphically in Figure 3.

Assumption 5 (Mediator ignorability). $Y_i(d, m) \perp\!\!\!\perp M_i(0) | D_i = d, M_i(1) = 1, X_i$

This is related to but dramatically stronger than Assumption 3, which merely requires that always-stop encounters are at least as severe in terms of observed criminal behavior. In contrast, for Assumption 5 to hold, violence rates in always-stop encounters must be *identical* to those in observationally equivalent racial stops. We find mediator ignorability to be highly implausible in the context of policing. Subjective factors such as an officer's suspicion and sense of threat—depicted as

U in Figure 3(c)—can not only lead to investigation (stopping) but also a heightened willingness to use force. These mediator-outcome confounders must be captured in *X* for this assumption to hold, but they are notoriously difficult to capture in officers' self-reported accounts. Even when proxies based on qualitative officer narratives are available, strong legal incentives exist for distortion. Moreover, analysts must be sure to condition on all variables related to officer mindset that are causally upstream of stops, while taking care not to induce bias by conditioning on any that are downstream.

Below, we demonstrate that every analysis estimating a racial effect using only data on stopped individuals implicitly makes Assumption 5. We further note that Assumptions 4(a), 4(b), and 5 are jointly covered by the slightly stronger assumption of sequential ignorability (Imai et al. 2011).

Assumption 6 (No Racial Stops). $M_i(0) = M_i(1)|M_i = 1$.

In Figure 3, this amounts to assuming away the arrow between *D* and *M*. Equivalently, this assumption states that all reported encounters were of the always-stop kind, or that there is no racial discrimination in stops. We show below that this assumption is implicitly made by all studies claiming to identify the average treatment effect of race, conditional on a reported interaction. Naturally, when there is no variation in $M_i(0)$, then this variable is ignorable and Assumption 5 is also satisfied.

However, in view of an overwhelming body of qualitative evidence and consistently massive quantitative differences in racial detainment rates across numerous policing domains, we find racial bias in police stops too plausible to dismiss by assumption (Alexander 2010; Baumgartner et al. 2017; Goel et al. 2016; Glaser 2014; Lerman and Weaver 2014). A raft of studies have also found that racial disparities persist even after leading candidate omitted variables, such as differential criminal activity across racial groups, are accounted for (Gelman et al. 2007). While such patterns are not proof of a causal relationship, we consider the possibility that police exhibit anti-minority bias when engaging civilians strong enough to merit a careful consideration of the implications of that bias for the validity of studies of racially biased policing.

Bias in the naïve estimator

In this section, we clear up several misunderstandings about the conventional estimator, which compares reported minority stops to reported white stops (with or without covariates). First, we show that when there is any racial discrimination in detainment, selection on stops introduces unavoidable statistical bias in estimating the $\text{ATE}_{M=1}$, even when a perfect set of observed covariates renders race ignorable with respect to the potential mediator and outcomes. These results directly contradict prior assertions that "linear regression can recover the 'race effect' if race is 'as good as randomly assigned,' conditional on the covariates" (Fryer 2018, 2). The issue is not one of omitted variables, but rather post-treatment conditioning. Second, we clarify an important open question about the nature of this bias. Fryer (2018) comments in the context of selection into arrest data that, "It is unclear how to estimate the extent of such bias or how to address it statistically," (5). Here, we derive the exact form of this bias for the $\text{ATE}_{M=1}$ and the $\text{ATT}_{M=1}$; Online Appendix A.3 does the same for the $\text{CDE}_{M=1}$. We show that the bias is always negative, resulting in naïve estimates that downplay the extent of racially discriminatory police violence. Below, we develop informative nonparametric sharp bounds that adjust the naïve estimates for the range of all possible selection bias.

Prior work on race and policing uses estimators that compare average reported outcomes in majority encounters to those in minority encounters. For simplicity of exposition, we present the special no-covariate case; Appendices A.1–A.3 derive the bias of the naïve estimator with covariate adjustment. We first refer readers to Equation 1, which expresses the naïve estimator $\hat{\Delta}$, in terms of stratum mean potential outcomes. We demonstrate that this commonly used analytic approach fails to recover any quantity of interest under plausible assumptions. We first show that it is biased for the $\text{ATE}_{M=1}$ and $\text{ATT}_{M=1}$ unless Assumption 6 is true and there are no racial stops. In Online Appendix A.3, we show it is also biased for the $\text{CDE}_{M=1}$ unless Assumption 5 holds—that is, always-stop encounters are identical in violence rates to racially discriminatory stops. As a result, the observed difference in means fails to recover any known causal quantity without additional, and highly implausible, assumptions.

In Online Appendix A.1, we derive the bias of $\hat{\Delta}$ when it is used to estimate $\text{ATE}_{M=1}$ under the

relatively plausible Assumptions 1–4. This bias can be written as:

$$\mathbb{E}[\hat{\Delta}] - \text{ATE}_{M=1}$$

$$= \big(\mathbb{E}[Y_i(1, 1) - Y_i(0, 1)|M_i(1) = 1, M_i(0) = 1]$$

$$- \mathbb{E}[Y_i(1, 1) - Y_i(0, 0)|M_i(1) = 1, M_i(0) = 0]$$

$$\big) \Pr(M_i(0) = 0|D_i = 1, M_i = 1)\Pr(D_i = 1|M_i = 1)$$

$$- \big(\mathbb{E}[Y_i(1, 1)|M_i(1) = 1, M_i(0) = 1]$$

$$- \mathbb{E}[Y_i(1, 1)|M_i(1) = 1, M_i(0) = 0]$$

$$\big) \Pr(M_i(0) = 0|D_i = 1, M_i = 1). \tag{6}$$

We offer several comments on Equation 6. The bias term is guaranteed to be negative, even with a perfect set of controls that render D_i ignorable, as long as there exist any racially discriminatory stops of minority civilians (or in an empirically falsified edge case).[8] The first term in the bias expression relates to heterogeneity in the average treatment effect, or the extent to which $Y_i(1, M_i(1)) - Y_i(0, M_i(0))$ differs in expectation between always-stop and racial-stop encounters—respectively, those with $M_i(1) = M_i(0) = 1$ and $M_i(0) < M_i(1)$.[9] Bias arises because in the latter type of encounter, a white civilian would never have been detained in the first place, and hence force would never have been used—that is, $\mathbb{E}[Y_i(0, 0)|D_i = 1, M_i(1) = 1, M_i(0) = 0] = 0$. Estimating the average potential outcomes of this group using stopped white civilians introduces unavoidable bias that the analyst cannot hope to eliminate simply by adding additional covariates to the estimating model. The second term is related to the difference in baseline violence rates between always-stop encounters and racially discriminatory stops; this term also vanishes if there are no racial stops.

Can the naïve estimator be rehabilitated by simply redefining the quantity of interest? In Online Appendices A.2–A.3, we show that the answer is no. The structure of the bias when $\hat{\Delta}$ is used

---

[8]The edge case is if there is zero use of force against white civilians. This possibility is empirically falsifiable; in our application, we show that it is far from the truth. To see that the bias is negative, observe that $\mathbb{E}[Y_i(0, 1)|M_i(1) = 1, M_i(0) = 1] \geq \mathbb{E}[Y_i(0, 0)|M_i(1) = 1, M_i(0) = 1]$, because the latter term is zero under Assumption 1. Together with Assumption 3, this signs the bias.

[9]Note that $M_i(d)$ simplifies in Equation 6, because it is constant within principal strata.

to estimate the $\text{ATT}_{M=1}$ is simpler, but leads to substantively identical conclusions: the naïve estimator is biased unless there are no racial stops. We show that bias for the $\text{ATT}_{M=1}$ is given by $\mathbb{E}[\hat{\Delta}] - \text{ATT}_{M=1} = -\mathbb{E}[Y_i(0,1)|M_i(1) = 1, M_i(0) = 1]\Pr(M_i(0) = 0|M_i(1) = 1)$. While the identifying assumptions for the $\text{CDE}_{M=1}$ are slightly weaker, they are nonetheless wholly implausible. The sign of this bias for the $\text{ATT}_{M=1}$ and $\text{CDE}_{M=1}$ can also be shown to be negative under Assumption 1–4, except in the implausible edge cases described in the Online Appendix. Thus, regardless of the target quantity, use of the observed difference in means will understate the rate of racially discriminatory police violence. In addition, we emphasize that these derivations show that statistical bias remains even after assuming a "complete" set of control variables that renders race ignorable. Post-treatment conditioning induces bias unless additional assumptions hold.

POTENTIAL SOLUTIONS

How should the analyst proceed in light of these results? We propose two approaches that eliminate the highly implausible assumptions outlined in the "Strong Assumptions" section, which are unstated but implicit in prior work. We caution that these solutions still rely on the weaker assumptions described in the "Necessary Assumptions" section, although we argue that these are often plausible in light of insights from extensive research on policing. Reasonable people can disagree on the plausibility of various assumptions, but by stating them explicitly, we seek to advance empirical work in an area which, at present, largely ignores such issues altogether.

In the first approach, we derive nonparametric sharp bounds representing the tightest possible range of causal effects that are consistent with the reported data (Manski 1995). Again, for simplicity, we begin by presenting bounds for the case in which treatment is unconditionally ignorable. To incorporate covariates, Online Appendix A.4 then describes a more general formulation in which bounds are computed within levels of X, without functional form assumptions, and reaggregated; this latter formulation is also applicable when a correctly specified regression is used. Both cases are demonstrated in a reanalysis of Fryer (2019) below.

A key limitation of the first proposed solution is that all quantities of interest remain only partially identified. This is fundamentally a consequence of selection into police administrative records; point

identification simply cannot be achieved without either implausible assumptions or additional data. To this end, we outline an alternative approach that incorporates limited information about the missing encounters (those that do not result in a stop). We show that with additional data—which in some cases are already being collected by agencies—the prevalence of racially discriminatory stops and most racial effects of interest can be point identified. Following our applied example, we describe a feasible research design based on this approach in detail.

Bounds on effect of race

Here, we derive large-sample nonparametric sharp bounds on the $\text{ATE}_{M=1}$ and $\text{ATT}_{M=1}$, focusing first on the case in which Assumption 4 (treatment ignorability) holds without conditioning on further covariates. Proposition 1 quantifies and corrects for the range of possible bias induced by post-treatment conditioning, producing an informative interval of possible joint values for (1) the partially identified $\text{ATE}_{M=1}$ and (2) the proportion of racial stops among reported minority encounters, $\rho = \Pr(M_i(0) = 0 | D_i = 1, M_i = 1)$. As Equation 6 suggests, when there is no racial bias in police stops ($\rho = 0$), these bounds collapse on the observed difference in means. We further demonstrate in Figure 4 that these bounds are highly informative when ρ is known or can be credibly estimated from supplemental data. When the prevalence of racially discriminatory detainment is unknown but a plausible range can be inferred from prior work, Figure 4 (discussed below) illustrates how this value can be used to assess the behavior of the bounds much like a sensitivity parameter.

Proposition 1 (Nonparametric Sharp Bounds on $\text{ATE}_{M=1}$)**.** *When D_i is ignorable, nonparametric sharp bounds on $(\text{ATE}_{M=1}, \rho)$ under Assumptions 1–4 are jointly given by*

$$\mathbb{E}[\hat{\Delta}] + \rho\, \mathbb{E}[Y_i | D_i = 0, M_i = 1]\,(1 - Pr(D_i = 0 | M_i = 1))$$

$$\leq \text{ATE}_{M=1} \leq$$

$$\mathbb{E}[\hat{\Delta}] + \frac{\rho}{1-\rho}\left(\mathbb{E}[Y_i | D_i = 1, M_i = 1] - \max\left\{0, 1 + \frac{1}{\rho}\mathbb{E}[Y_i | D_i = 1, M_i = 1] - \frac{1}{\rho}\right\}\right) Pr(D_i = 0 | M_i = 1)$$

$$+\rho\, \mathbb{E}[Y_i | D_i = 0, M_i = 1]\,(1 - Pr(D_i = 0 | M_i = 1)).$$

where $\hat{\Delta} = \overline{Y_i | D_i = 1, M_i = 1} - \overline{Y_i | D_i = 0, M_i = 1}$, *and the* $(ATT_{M=1}, \rho)$ *must similarly satisfy*

$$ATT_{M=1} = \mathbb{E}[\hat{\Delta}] + \rho\,\mathbb{E}[Y_i | D_i = 0, M_i = 1]$$

To derive Proposition 1, we reformulate the bias in terms of the unobserved joint distribution of (1) use of force in minority encounters and (2) whether a minority stop was racially discriminatory. Following Knox et al. (2019), we then use Assumptions 1–4 and the Fréchet inequalities, in conjunction with the observed margins, to place sharp bounds on this joint distribution. These then imply sharp bounds on the $\text{ATE}_{M=1}$. A detailed proof is given in Online Appendix A.4 for the more general case in which D_i is ignorable only after conditioning on pre-stop covariates. In this case, the local average treatment effect, $\text{ATE}_{M=1,x}$, is first bounded by applying Proposition 1 within levels of X to obtain local bounds, $[\underline{\text{ATE}}_{M=1,x}, \overline{\text{ATE}}_{M=1,x}]$. These are then straightforwardly reaggregated to obtain bounds on the conditional treatment effect among stops, $[\sum_x \underline{\text{ATE}}_{M=1,x}\, \Pr(X_i = x | M_i = 1)\,,$ $\sum_x \overline{\text{ATE}}_{M=1,x}\, \Pr(X_i = x | M_i = 1)]$. In Online Appendix A.5, we outline a Monte Carlo procedure for constructing confidence intervals that asymptotically contain both the true lower and upper bounds endpoints with probability $1 - \alpha$.

We note that the proportion of racially discriminatory stops may vary with X. However, when using these bounds as a sensitivity analysis, we suggest using the simplifying approximation of a constant ρ. This is because without additional data beyond civilian race, use of force, or even pre-stop covariates, police administrative records alone are virtually uninformative about the range of ρ: any value in $[0, 1)$ could produce the observed data,[10] although Proposition 1 shows that each possible ρ value has differing implications for the set of possible racial effects.

Point identification of the ATE given additional data

The ATE is point identified with the collection of only two additional numbers—the count of total minority and white encounters, within levels of X where applicable. Below, we propose an alternative design in which these data are collected from passive instruments such as traffic cameras or police

---

[10]If all stops were racially discriminatory, then we would observe no white stops.

body-worn cameras. Where such a design is infeasible (for example, where traffic cameras cover only a subset of the jurisdiction under study), point identification can also be achieved by linking incomplete data on both reported and unreported encounters to police administrative records under mild assumptions.

Proposition 2 (Point Identification of ATE). *Under Assumptions 1–4, the ATE is identified by a weighted combination of the observed racial means,*

$$\mathbb{E}[Y_i|D_i = 1, M_i(D_i) = 1] \, Pr(M_i = 1|D_i = 1) - \mathbb{E}[Y_i|D_i = 0, M_i(D_i) = 1] \, Pr(M_i = 1|D_i = 0).$$

Intuitively, the proof breaks the ATE into the size-weighted sum of principal effects among always-stop and racial-stop encounters (the principal effect in never-stop encounters is known to be zero). Crucially, the additional data on non-stops allows the researcher to construct a contingency table representing the joint distribution of race and detainment. As part of the proof in Online Appendix A.6, we show that this can be used to straightforwardly recover the size of each principal stratum under Assumptions 2 and 4(a). However, it remains impossible to determine whether any individual stop was racially discriminatory.

When total encounter numbers are unknown, this joint distribution can nonetheless be estimated by attempting to link a representative sample of all encounters (e.g. using timestamps from traffic cameras) against administrative records (e.g. license plate databases); those that are unlinkable can be presumed unreported. After recovering principal strata sizes, we then proceed by noting that minority outcomes in reported administrative data are in fact a mixture of $Y_i(1, M_i(1))$ from both always-stop and racial-stop strata in precisely the required proportions; that reported white outcomes correspond to $Y_i(0, M_i(0))$ from the always-stop stratum; and that $Y_i(0, M_i(0))$ is known to be zero among the racial-stop stratum under Assumption 1. From this, the ATE can then be reconstructed.

REANALYSIS OF FRYER (2019)

We have shown that the standard approach to estimating racial bias in police data will always underestimate its degree, so long as police discriminate against minorities when choosing whom to

investigate. To explore the magnitude of this statistical bias in an applied setting, we replicate and extend a section of Fryer (2019) which reports estimates of racial discrimination in the application of sub-lethal force using the NYPD's "Stop, Question and Frisk" (SQF) database (2003-2013).[11] The NYPD data contain roughly 5 million records of pedestrian stops, the vast majority of which are of nonwhite suspects. The data record the use of varying levels of force, including laying hands on a suspect, handcuffing a suspect, pointing a weapon at a suspect, and pepper spraying a suspect, among others. The original analysis in Fryer (2019) utilized the simple naïve approach of Equation 1 to predict the severity of force applied by police, as well as covariate-adjusted naïve models analogous to those we consider in Appendices A.1–A.3. Specifically, the study presented a logistic regression of police force on suspect race, along with additional specifications that added a host of control variables such as precinct fixed effects, to render the ignorability assumptions more plausible. We reproduce two of these models—the baseline specification including only racial group indicators, along with the richer "main specification" (21)[12]—to estimate the conditional expectations in Proposition 1. For comparability to the original analysis, we take these models at face value, setting aside issues of potential model misspecification and the ignorability of civilian race.

One analysis in Fryer (2019) considered the use of any force against a suspect, while subsequent analyses examined force exceeding various severity thresholds, such as a binary outcome for "at least use of handcuffs." Using the coding rules and estimation procedures in Fryer (2019), we were able to closely replicate the published results. However, in doing so, we discovered this procedure involved an unconventional and inadvisable step in which all observations with non-zero force below

---

[11]Because the replication material for Fryer (2019) was not posted at the time of analysis, these data were obtained directly from `https://www1.nyc.gov/site/nypd/stats/reports-analysis/stopfrisk.page`.

[12]The main specification in Fryer (2019) consists of a logistic regression of a force outcome on race dummies plus controls for gender, a quadratic in age, whether the stop was indoors or outdoors, whether the stop took place during the daytime, whether the stop took place in a high crime area, during a high crime time, or in a high crime area at a high crime time, whether the officer was in uniform, civilian ID type, whether others were stopped during the interaction, controls for civilian behavior, and precinct and year fixed effects. Fryer (2019) also notes that "missing indicators for all variables" are included as covariates. We omit these indicators as it was unclear how they were coded. See Figure 1 caption in Fryer (2019).

the threshold of interest were dropped—a severe case of selection on the dependent variable. In the most extreme case, in the analysis of police baton and pepper spray use, this resulted in the discarding of all encounters in which only lower levels of force were used, a set that comprised 21.5% of all observations and 99.8% of all uses of force. In order to present the most defensible results possible, for these outcomes, we depart from the analysis in Fryer (2019) and revise the procedure so that all encounters with a level of force at or above a given threshold are assigned an outcome of 1 (as before) and all other encounters are assigned a value of 0 (including those with lower levels of force, which are now retained). Section B.1 in the Online Appendix contains an extended discussion of the issue; a comparison of the original, replicated, and corrected results; and a demonstration of the serious implications for statistical significance of the original estimates.

Based on the discussion in both Fryer (2019) and Fryer (2018), we interpret the published results as estimates of the $\text{ATE}_{M=1}$: "the difference in Y that can be attributed to an individual's race," (Fryer 2018, 2), conditional on a recorded interaction with police (that is, conditional on $M_i = 1$). We note that of the other quantities considered in this paper, the unconditional ATE cannot be estimated without information on unreported encounters, and the $\text{CDE}_{M=1}$ cannot be computed without strong assumptions about potential outcomes that can never be realized in observational settings. For these reasons, we focus on the $\text{ATE}_{M=1}$ and $\text{ATT}_{M=1}$ in this reanalysis.[13]

---

[13]We note that in Proposition 1 we consider binary minority status, whereas the specifications in Fryer (2019) take civilian race as a categorical variable. (However, only two races are considered for any particular $\text{ATE}_{M=1}$ estimate: black versus white, or Hispanic versus white). To accommodate this, in reported black $\text{ATE}_{M=1}$ and $\text{ATT}_{M=1}$ results, we use a slight generalization in which white civilian encounters are represented with $D_i = 0$, black encounters with $D_i = 1$, and subsequent minority groups with $D_i = 2, 3$ and so on. Proposition 1 and its covariate-adjusted counterpart in Online Appendix A.4 can then be applied directly. The chief implication of this formulation is (1) a different average value for $Y_i(d, 1)$ is estimated for each minority group, and (2) that all minority groups are implicitly assumed to be racially stopped at the same rate, although this can easily be relaxed. (The same procedure is applied when the minority group of interest is Hispanic civilians, after setting the Hispanic indicator to $D_i = 1$.) To assess whether results were affected by this, in Online Appendix B.4, we conduct two additional analyses after first subsetting to black and white encounters, and Hispanic and white encounters, respectively. As the results makes clear, conclusions are virtually identical apart from differences

Figure 4 depicts bounds on the $\text{ATE}_{M=1}$ when the binary outcome is any use of force, including the lowest recorded value of physically handling a civilian.[14] Importantly, this specific outcome is unaffected by the outcome coding issue discussed above. (In Figures B2 and B3, we present additional bounds for varying force thresholds, up to whether a baton or pepper spray was used.) The plots also display estimates of the bias-corrected $\text{ATT}_{M=1}$ (dashed lines). As the plots show, the range of possible $\text{ATE}_{M=1}$ and $\text{ATT}_{M=1}$ values varies strongly with the severity of discrimination in stops.

In Equation 6, we demonstrated that use of the naïve estimator implied the substantively implausible assumption that police never discriminate in stops (i.e. $\rho = 0$). Similarly, contextual information also suggests that some depicted values of ρ are implausibly large. To understand the range of empirically plausible values, we turn to two prior studies that use very different analytic approaches to shed light on the degree of racial bias in the decision to detain civilians. Using the SQF data and controlling for precinct, suspected crime, and prior local arrest rates by race, Gelman et al. (2007) produce estimates that—by our calculations—imply 32% of black-civilian stops made by the NYPD could not be explained even by differential criminality between racial groups of suspects, as proxied by prior arrest rates.[15] Their analyses are run separately by precinct and crime type; for simplicity, we take the weighted average of racial-stop proportions. This analytic approach most likely underestimates the proportion of racially discriminatory stops—the number of prior arrests in a precinct and racial group is not a direct measure of criminality, but is itself contaminated by discrimination in previous detainments

---

that stem from the size of the subsetted data.

[14]Note that we treat stops in which "other" was denoted as the use of force category as zero force, since the vast majority of these cases did not even involve officers laying hands on suspects.

[15]Based on SQF data from 1998-1999, Gelman et al. (2007) fit hierarchical Poisson models for the number of stops (by suspected crime, precinct, and race) per arrest in the previous year, which they model as $e^{\mu + \alpha_{\text{race}}}$ within groups of stops defined by the suspected charges (violent crimes, weapons crimes, property crimes, and drug crimes) and precinct racial composition (<10%, 10–40%, and >40% black). Within each group, the excess black stopping rate is then given by $1 - e^{\alpha_{\text{white}} - \alpha_{\text{black}}}$. We approximate the size of each group by multiplying the reported marginal probabilities of stop types (25%, 44%, 20%, and 11%, respectively) and composition groups ("each... represents roughly 1/3 of the precincts"), since the joint distribution is not reported. The $\rho = 0.32$ estimate is then produced by taking the size-weighted average of subgroup excess black stopping rates. The corresponding estimate of ρ for Hispanic civilians implied by Gelman et al. (2007) is slightly higher, at 0.35.

and arrests. We therefore regard the value of ρ implied by Gelman et al. (2007) as conservative.

Goel et al. (2016) take an entirely different tack based on a comparison of "hit rates," or the share of stops that produced evidence of the suspected crime for which the civilian was detained—a variant of an "outcome test" for discrimination (Anwar and Fang 2006; Knowles et al. 2001). Using a flexible logistic regression to adjust for a vast array of indicators visible to officers pre-stop, the study shows that white hit rates exceeded those of "similarly situated" black civilians. We show in our Online Appendix (section A.7) that the difference in hit rates implies a minimum proportion of racial stops, and therefore also implies a conservative estimate of ρ.[16] The corresponding values of ρ from these two studies are 0.32 and a lower bound of 0.34, respectively, when considering black civilians. While any estimate of this difficult-to-measure quantity from police data is sure to be imperfect, the fact that two independent estimates of racial bias in stopping so closely comport with one another, despite using wholly different analytical approaches, gives us some empirical justification for narrowing the range of plausible racial effects in the use-of-force analysis. We note that the research design presented in the "Recommendations for Future Research" section below offers an alternative approach for obtaining better estimates of racially discriminatory stopping.

Figure 4 demonstrates that strong negative bias in the naïve estimator paints a wildly misleading portrait of police use of force. We turn first to estimates of the $\text{ATE}_{M=1}$ using the main specification, which adjusts for a battery of covariates. The naïve estimator (which assumes no racial bias in police stops) suggests that encounters with black (Hispanic) suspects are predicted to exhibit an additional 3.9 (0.4) instances of handcuffing per 1,000 encounters, compared with the same encounters had they involved white civilians. We then employ the most conservative racial stopping estimate, denoted by the vertical line in the figure, to generate bounds on the true race effect. Our bias-corrected results show the true effect is at least as high as 15.5 (13.0)—meaning that the conventional approach underestimates discriminatory force by a factor of at least 4 (32).

To characterize bias in estimates of the $\text{ATT}_{M=1}$, we again use the conservative racial stopping estimate from Gelman et al. (2007) to correct the naïve estimate. Again, the naïve approach substantially

---

[16]Using SQF data from 2008–2012, Goel et al. (2016) estimate a hit rate of 3.8% for white suspects and 2.5% for black suspects (379), which implies that ρ is at least 0.34.

FIGURE 4. Bounds for racially discriminatory use of force, any severity. These plots present the $\text{ATE}_{M=1}$ ($\text{ATT}_{M=1}$) for excess racial force, scaled by the number of stops (number of minority stops) to obtain the total number of civilians affected. The left panels consider the difference in use of force if black civilians were substituted into each encounter of any race (each black encounter), versus white civilians; the right panels show the same quantities for Hispanic civilians. Blue points (error bars) denote the naïve estimator (95% confidence intervals), which, conditional on the typical selection-on-observables assumption, is unbiased for the $\text{ATE}_{M=1}$ if there are no discriminatory stops of minority civilians (zero on the x-axis). The dark (light) regions represent the range of possible values (95% CI) for (1) the $\text{ATE}_{M=1}$ and (2) the proportion of discriminatory stops in reported data jointly, per Proposition 1. The vertical line corresponds to an estimate of the proportion of discriminatory stops from Gelman et al. (2007), suggesting a plausible value for this unobservable parameter. The top (bottom) panels present bounds based on a model with no controls (the main specification, adjusting for a wide range of covariates).
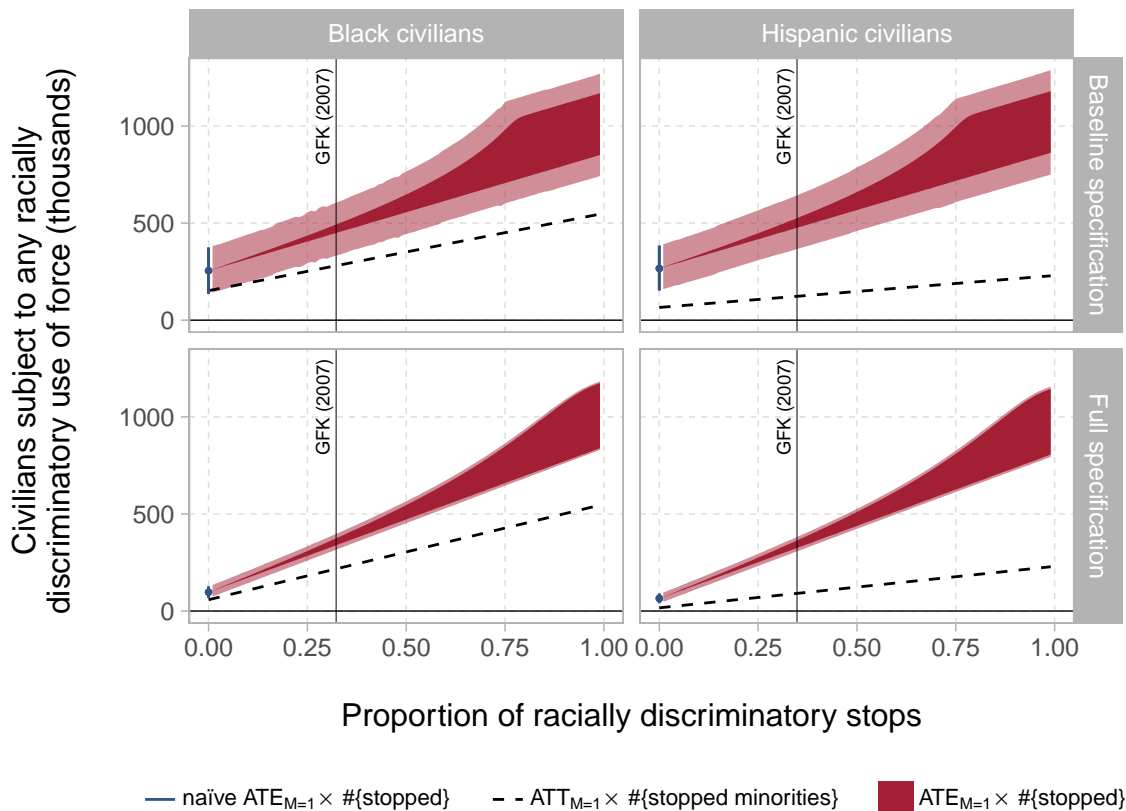
understates racially discriminatory police violence, suggesting that there were 75,000 instances in which police laid hands on black and Hispanic civilians, but would not have done so had those individuals been white. Our bias-corrected estimate shows the true number is approximately 307,000, meaning the naïve approach masks 232,000 such incidents. Similarly, the naïve approach indicates roughly 3,400 racially discriminatory instances in which officers pointed a weapon at a black or Hispanic civilian, whereas the bias-corrected $\text{ATT}_{M=1}$ shows the true number is almost five times as large.

To see how this statistical bias affects estimates for different levels of force, Table 1 presents naïve estimates alongside $\text{ATE}_{M=1}$ bounds for excess force per 1,000 black and Hispanic encounters across the full spectrum of police actions—ranging from physical handling of a civilian to the use of pepper spray or a baton—again using the conservative racial-stop estimate from Gelman et al. (2007) to apply our bias correction. The results again show that the traditional approach substantially understates the degree of racial bias in police use of force. Our results also include numerous cases in which downward bias produces the illusion of no race effect. For example, while the approach in Fryer (2019) implies a statistically insignificant 2.4 instances per 1,000 encounters of pushing Hispanic suspects to a wall due to suspect race, our revised estimate shows the true number is at least 26—eleven times as large. We can also quantify the number of masked instances of racially discriminatory uses of force as a percentage of all uses of force, displayed in Figure 5. In the period we examine, black and Hispanic civilians experience force at the hands of police 779,894 times. Using the approach in (Fryer 2019), one would conclude that about 10% would not have occurred had those civilians been white. Using our bias-corrected approach, we find that in fact 39% were discriminatory. These underestimates persist across all force threshold analyses.[17]

---

[17]These estimates were generated by computing the $\text{ATE}_{M=1}$ with covariate adjustment.

FIGURE 5. Estimated number of racially discriminatory uses of force against black and His-panic civilians, divided by total observed uses of force among those groups using naïve (red dot) and bias-corrected (blue triangle) estimators of the $ATT_{M=1}$. In some cases, the naïve ap-proach returns negative estimates, indicating that *more* uses of force would have occurred had the civilians been white. The bias-corrected estimates show the naïve estimates substantially underestimate the pervasiveness of racial bias in police violence.

TABLE 1. Average treatment effect among stops (ATE$_{M=1}$), by severity of force and minority group. Excess use of force used against minority civilians (versus white civilians) per 1,000 encounters. Bounds intervals indicate the range of possible ATE$_{M=1}$ values when the unknown proportion of discriminatory stops is approximated with the conservative estimate from Gelman et al. (2007). Estimates are bolded, and 95% confidence intervals are italicized.

| | ATE$_{M=1}$ for encounters with black civilians (vs. white) | | | |
| | No covariates | | Full specification | |
Minimum force	bounds	naïve	bounds	naïve
Use of hands	**(112.66, 124.59)**	**61.69**	**(86.95, 96.70)**	**23.46**
	(84.6, 151.84)	*(32.89, 90.63)*	*(81.54, 102.14)*	*(16.23, 30.54)*
Push to wall	**(24.15, 27.75)**	**4.20**	**(26.47, 30.20)**	**6.66**
	(15.50, 37.35)	*(-5.29, 14.02)*	*(24.24, 32.39)*	*(3.67, 9.53)*
Use of handcuffs	**(14.60, 16.92)**	**1.32**	**(16.59, 19.05)**	**3.95**
	(9.45, 22.61)	*(-4.83, 7.53)*	*(15.10, 20.57)*	*(1.95, 5.90)*
Draw weapon	**(4.52, 5.14)**	**1.26**	**(4.71, 5.34)**	**1.46**
	(3.13, 6.67)	*(-0.33, 2.83)*	*(4.20, 5.86)*	*(0.77, 2.14)*
Push to ground	**(4.04, 4.58)**	**1.22**	**(4.10, 4.64)**	**1.24**
	(2.79, 5.97)	*(-0.21, 2.66)*	*(3.65, 5.07)*	*(0.64, 1.80)*
Point weapon	**(1.49, 1.70)**	**0.36**	**(1.63, 1.86)**	**0.55**
	(0.96, 2.29)	*(-0.29, 1.00)*	*(1.36, 2.12)*	*(0.18, 0.89)*
Baton or pepper spray	**(0.17, 0.19)**	**0.08**	**(0.17, 0.19)**	**0.07**
	(0.10, 0.26)	*(-0.01, 0.15)*	*(0.12, 0.24)*	*(0.00, 0.14)*

| | ATE$_{M=1}$ for encounters with Hispanic civilians (vs. white) | | | |
| | No covariates | | Full specification | |
	bounds	naïve	bounds	naïve
Use of hands	**(115.44, 127.53)**	**64.48**	**(78.93, 88.31)**	**15.44**
	(88.94, 155.96)	*(37.06, 92.91)*	*(74.70, 92.65)*	*(9.60, 21.09)*
Push to wall	**(26.41, 30.14)**	**6.46**	**(22.24, 25.75)**	**2.44**
	(19.54, 37.79)	*(-1.12, 14.26)*	*(20.21, 27.80)*	*(-0.28, 5.08)*
Use of handcuffs	**(12.54, 14.76)**	**-0.74**	**(13.05, 15.31)**	**0.40**
	(9.10, 18.24)	*(-5.27, 3.57)*	*(11.71, 16.64)*	*(-1.40, 2.1)*
Draw weapon	**(3.42, 3.98)**	**0.16**	**(3.12, 3.66)**	**-0.14**
	(2.41, 5.08)	*(-1.04, 1.33)*	*(2.63, 4.16)*	*(-0.77, 0.49)*
Push to ground	**(3.11, 3.60)**	**0.29**	**(2.71, 3.18)**	**-0.14**
	(2.18, 4.61)	*(-0.83, 1.37)*	*(2.27, 3.63)*	*(-0.71, 0.41)*
Point weapon	**(0.73, 0.90)**	**-0.41**	**(0.81, 0.98)**	**-0.28**
	(0.32, 1.29)	*(-0.94, 0.08)*	*(0.54, 1.26)*	*(-0.64, 0.07)*
Baton or pepper spray	**(0.05, 0.06)**	**-0.05**	**(0.05, 0.07)**	**-0.05**
	(-0.01, 0.12)	*(-0.13, 0.02)*	*(0.00, 0.12)*	*(-0.12, 0.02)*

RECOMMENDATIONS FOR FUTURE RESEARCH

The analysis above clarifies whether and when estimates of racial bias in police behavior identify causal quantities, shedding light on how traditional estimation approaches that fail to account for post-treatment conditioning can inadvertently mask racially biased policing. Our results suggest the body of evidence on this topic that relies on police administrative data may be largely uninformative or even misleading. While our bias-correction and bounding techniques are an improvement, they still rely on assumptions that many analysts may not be willing to entertain. Some of these assumptions, such as conditional treatment ignorability, are unavoidable. But others can be sidestepped or weakened through the use of research designs that preempt the problem of post-treatment conditioning. In what follows, we detail a feasible research design that addresses these concerns.

To estimate the effect of suspect race on post-stop police behavior while avoiding the concerns outlined above, we describe a feasible study of police-civilian interactions during traffic stops. A key advantage of traffic studies is that much of the data needed to improve research are already collected passively by law enforcement agencies across the U.S. in an automated fashion via highway cameras. We note that before the advent of this technology, data on unreported police-civilian interactions had to be manually collected by researchers accompanying patrol officers on their shifts (Allen 1982; Smith et al. 1984), a labor-intensive strategy highly vulnerable to researcher demand effects (Orne 1962).

Recall that a key problem in the typical study of police administrative data is the unobservability of those encounters that do not generate police reports. However, given the prevalence of highway speed cameras across police jurisdictions, it is entirely feasible to collect data on every passing car (or a random sample of passing cars), whether or not police pulled the car over and recorded the stop. This mode of data collection has already been utilized in prior work (Kocieniewski 2002; Lange et al. 2005), though in those studies camera data on individual motorists were not linked to administrative data on policing outcomes, as we propose below.

Given a large random sample of passing cars captured by highway speed cameras, analysts could use video or photographic records to document license plate numbers that allow for a merge with other administrative data sets containing information on the registrant's home neighborhood, whether each car went on to be stopped by nearby police at a proximate time, whether a summons was issued, and

whether the encounter escalated to include a search or the use of force. As with all causal analyses of observational data, analysts must still make some version of Assumption 4(b)—no treatment-outcome confounding conditional on observable covariates—but in this case, the standard "treatment selection on observables" plausibly holds, because virtually all pre-stop data available to an officer are in fact literally observable to the analyst. Using camera footage merged with administrative records, analysts could credibly measure this "complete" set of control variables.[18] These factors would include not only the race, age, gender, and registered neighborhood of the driver, but also the make, color, and condition of the car, along with weather and driving speed.

Given this set of covariates, researchers could credibly estimate the ATE for various outcomes, including searching, ticketing and the use of force, by comparing the rates of outcomes between racial minority and majority motorists, regardless of whether they were stopped by police, conditional on X. The $\text{ATT}_{M=1}$ is similarly point identified, because the proportion of racial stops can be calculated and used to correct estimates. However, the $\text{ATE}_{M=1}$ remains partially identified—the quantity can be bounded, as we show above, but not precisely estimated. And as Figure 6 makes clear, the $\text{CDE}_{M=1}$ remains fundamentally unidentifiable without covariates that make Assumption 5 plausible, such as controls for officer temperament that are specific to some stops but not others (i.e., time-varying), which likely influences both stopping decisions and subsequent treatment of civilians.

CONCLUSION

With the release of large and granular data on police-civilian interactions, many researchers have focused on estimating whether police exhibit racial bias in their treatment of civilians. Though some studies have acknowledged the threat of post-treatment bias in this setting (Fryer 2018), the issue has not been adequately addressed, and studies in this area have left ambiguous which causal quantities

---

[18]This approach is akin to the design of Hainmueller and Hangartner (2013), another rare instance in which the analyst could claim to measure all relevant covariates in an observational setting. In that study, citizens made judgments about individuals applying for citizenship in Switzerland. Because all information on potential citizens was contained on a flier distributed by the government, the authors could credibly account for all possible factors that contributed to the average citizen's judgment of applicants.

FIGURE 6. Traffic stop design. The DAG illustrates potential back-door paths for stops (through W, e.g., heavily policed neighborhoods) and for the use of force (through V, e.g., car registrant has warrant for arrest) that may correlate with the presence of minority drivers. These are blocked (boxed) by conditioning on pre-stop variables, including license plates as well as administrative records that can be linked through them. Many mediator-outcome confounders (U) cannot be blocked but do not pose a threat to inference for the ATE or $\text{ATE}_{M=1}$.

are being approximated and the degree to which racial bias may be obscured by traditional estimation strategies. Given the policy relevance of this topic and the degree of selection bias inherent to these analyses, we believe social scientists need to devote substantial effort to develop research designs that can sidestep the threat of post-treatment conditioning rather than proceeding in the face of this threat and simply hoping for the best.

In this paper, we clarify the statistical problems in the use of police administrative data in isolation to study racial bias. We offer bias-correction and bounding procedures for scholars analyzing these data, along with an improved research design that can avoid post-treatment conditioning altogether. Our results can inform the study of racial discrimination in a host of other settings beyond law enforcement. And though we focus on a case of racial bias in the U.S., these results also speak to a rich literature on racial discrimination outside the U.S. context (e.g. Bruce-Jones 2015; Cano 2010). Our identifying assumptions may also be useful for researchers seeking to address biases stemming from post-treatment conditioning more generally, beyond studies of discrimination.

While we are optimistic about alternative designs and estimation strategies, we are under no illusions that eliminating this particular source of bias will remove others. Our research design suggestions may also limit the outcomes that are feasible to study. For example, rare events like shootings may or may not occur during the observation periods proposed, meaning only lower-level uses of force or sanctioning can be studied in some cases. Our recommendations therefore place emphasis on bias reduction over latitude in the selection of research questions. But given the ease with which faulty conclusions can be reached as a result of the race-based selection we highlight, narrowing the scope of

research in order to generate more reliable estimates may be preferable, especially since policy reforms could hinge on the results of studies in this area. Put differently, because of the pitfalls we highlight above, it is not clear that studies of rare phenomena that lack a sound design are generating usable knowledge anyway, so this tradeoff in scope may be of only marginal concern (Samii 2016).

Regardless of which approach scholars pursue, this paper highlights the need for further careful research into the first stage of police-civilian interactions—that is, the process by which officers decide whether or not to stop and investigate an individual for a crime. This effort is necessary not only to further our scholarly understanding of police-civilian interactions, but to craft effective policy reforms. If racial bias is concentrated in the initial stage of contact, reforms focused on reducing unnecessary police-civilian interactions may be most effective at curbing racially-discriminatory police violence. On the other hand, if there exists more significant bias in the ultimate decision to use force, substantial improvements may require a wholly different reform strategy. Without serious consideration of the role of race in each stage of the complex police-civilian interactions under study, the benefits of data-driven reforms will be stunted, as will our collective understanding of the politics of policing.

REFERENCES

Acharya, Avidit , Matthew Blackwell, and Maya Sen (2016). Explaining causal findings without bias: Detecting and assessing direct effects. *Biometrics 110*(3), 512–529.

Alesina, Alberto and Eliana La Ferrara (2014). A test of racial bias in capital sentencing. *American Economic Review 104*(11), 3397–3433.

Alexander, Michelle (2010). *The New Jim Crow: Mass Incarceration in the Age of Colorblindness*. The New Press.

Allen, David (1982). Police supervision on the street: An analysis of supervisor/officer interaction during the shift. *Journal of Criminal Justice 10*(2), 91–109.

Angrist, Joshua D. , Guido W. Imbens, and Donald B. Rubin (1996). Identification of causal effects using instrumental variables. *Journal of the American Statistical Association 91*(434), 444–455.

Angrist, Joshua D. and Jörn-Steffen Pischke (2008). *Mostly Harmless Econometrics: An Empiricist's Companion.* Princeton University Press.

Antonovics, Kate and Brian G. Knight (2009). A new look at racial profiling: Evidence from the boston police department. *The Review of Economics and Statistics 91*(1), 163–177.

Anwar, Shamena and Hanming Fang (2006). An alternative test of racial prejudice in motor vehicle searches: Theory and evidence. *The Review of Economic Studies 96*(1), 127–151.

Arnold, D. , W. Dobbie, and C.S. Yang (2018). Racial bias in bail decisions. *The Quarterly Journal of Economics 133*(4), 1885–1932.

Arrow, Kenneth J. (1972). Models of job discrimination. In A. Pascal (Ed.), *Racial Discrimination in Economic Life*, pp. 83–102. D.C. Heath.

Arrow, Kenneth J. (1998). What has economics to say about racial discrimination? *Journal of economic perspectives 12*(2), 91–100.

Atiba Goff, Phillip and Kimberly Barsamian Kahn (2012). Racial bias in policing: Why we know less than we should. *Social Issues and Policy Review 6*(1), 177–210.

Ayres, Ian (2002). Outcome tests of racial disparities in police practices. *Justice Research and Policy 4*(1-2), 131–142.

Balke, Alexander and Judea Pearl (1997). Bounds on treatment effects from studies with imperfect compliance. *Journal of the American Statistical Association 92*(439), 1171–1176.

Baumgartner, Frank R. , Derek A. Epp, Kelsey Shoub, and Bayard Love (2017). Targeting young men of color for search and arrest during traffic stops: evidence from north carolina, 2002-2013. *Politics, Groups, and Identities 5*(1), 107–131.

Becker, Gary (1971). *The Economics of Discrimination.* University of Chicago Press.

Blackwell, Matthew (2013). A framework for dynamic causal inference in political science. *American Journal of Political Science 57*(2), 504–520.

Brehm, John and Scott Gates (1999). *Working, Shirking, and Sabotage: Bureaucratic Response to a Democratic Public.* University of Michigan Press.

Bruce-Jones, Eddie (2015). German policing at the intersection: race, gender, migrant status and mental health. *Race & Class 56*(3), 36–49.

Burch, Traci (2013). *Trading Democracy for Justice: Criminal Convictions and the Decline of Neighborhood Political Participation.* University of Chicago Press.

Cano, Ignácio (2010). Racial bias in police use of lethal force in brazil. *Police Practice and Research: An International Journal 11*(1), 31–43.

Cohen, Elisha , Anna Gunderson, Kaylyn Jackson, Paul Zachary, Tom S. Clark, Adam N. Glynn, and Michael Leo Owens (2017). Do officer-involved shootings reduce citizen contact with government? *The Journal of Politics 81*(3), 1111–1123.

Eberhardt, Jennifer , Phillip Atiba Goff, Valerie J. Purdie, and Paul G. Davies (2004). Seeing black: Race, crime, and visual processing. *Journal of Personality and Social Psychology 87*(6), 876–893.

Edwards, Frank , Hedwig Lee, and Michael Esposito (2019). Risk of being killed by police use of force in the united states by age, race–ethnicity, and sex. *Proceedings of the National Academy of Sciences 116*(34), 16793–16798.

Elwert, Felix and Christopher Winship (2014). Endogenous selection bias: The problem of conditioning on a collider variable. *The Annual Review of Sociology 40*, 31–53.

Fisher, Marc and Peter Hermann (2015, June 8). Did the mckinney, texas, police officer know he was being recorded? *The Washington Post*.

Frangakis, Constantine E. and Donald B. Rubin (2002). Principal stratification in causal inference. *Biometrics 58*(1), 21–29.

Fridell, Lorie A. (2017). Explaining the disparity in results across studies assessing racial disparity in police use of force: A research note. *Journal of economic perspectives 42*(3), 502–513.

Fryer, Roland G. (2018). Reconciling results on racial differences in police shootings. *American Economic Review (Papers and Proceedings) 108*, 228–233.

Fryer, Roland G. (2019). An empirical analysis of racial differences in police use of force. *Journal of Political Economy 127*(3), 1210–1261.

Gelman, Andrew , Jeffrey Fagan, and Alex Kiss (2007). An analysis of the new york city police department's "stop-and-frisk'" policy in the context of claims of racial bias. *Journal of the American Statistical Association 102*(429), 813–823.

Glaser, Jack (2014). *Suspect Race: Causes and Consequences of Racial Profiling*. Oxford University Press.

Goel, Sharad , Justin M. Rao, and Ravi Shroff (2016). Precinct or prejudice? understanding racial disparities in new york city's stop-and-frisk policy. *Annals of Applied Statistics 10*(1), 365–394.

Gottschalk, Marie (2008). Hiding in plain sight. *Annual Review of Political Science 11*(1), 235–260.

Greiner, James D. and Donald B. Rubin (2011). Causal effects of perceived immutable characteristics. *Review of Economics and Statistics 93*(4), 775–785.

Grogger, Jeffrey and Greg Ridgeway (2006). Testing for racial profiling in traffic stops from behind a veil of darkness. *Journal of the American Statistical Association 101*(475), 878–887.

Hainmueller, Jens and Dominik Hangartner (2013). Who gets a swiss passport? a natural experiment in immigrant discrimination. *American Political Science Review 107*(1), 159–187.

Harvey, Anna and Murat Mungan (2019). Policing for profit: The political economy of law enforcement. *Working Paper*. `https://s18798.pcdn.co/annaharvey/wp-content/uploads/sites/6417/2019/06/Policing_For_Profit_2019.pdf`.

Heckman, James J. (1979). Sample selection bias as a specification error. *Econometrica 47*(1), 153–161.

Hernán, Miguel , Sonia Hernández-Diáz, and James Robins (2004). A structural approach to selection bias. *Epidemiology 15*(5), 615–625.

Hernán, Miguel A. (2016). Does water kill? a call for less casual causal inferences. *Annals of Epidemiology 26*(10), 674–680.

Holland, Paul W. (1986). Statistics and causal inference. *Journal of the American Statistical Association 81*(396), 945–60.

Imai, Kosuke , Luke Keele, Dustin Tingley, and Teppei Yamamoto (2011). Unpacking the black box of causality: Learning about causal mechanisms from experimental and observational studies. *American Political Science Review 105*(4), 765–789.

Iyengar, Shanto (1994). *Is anyone responsible?: How television frames political issues*. University of Chicago Press.

Johnson, David J. , Trevor Tress, Nicole Burkel, Carley Taylor, and Joseph Cesario (2019). Officer characteristics and racial disparities in fatal officer-involved shootings. *Proceedings of the National Academy of Sciences 116*(32), 15877–15882. Published online ahead of print July 22, 2019. `https://www.pnas.org/content/early/2019/07/16/1903856116`.

Key, V.0. (1949). *Southern Politics in State and Nation*. Knopf.

Knowles, J. , N. Perisco, and P. Todd (2001). Racial bias in motor vehicle searches: Theory and evidence. *Journal of Political Economy 109*(1), 203–229.

Knox, Dean and Jonathan Mummolo (forthcoming). Making inferences about racial disparities in police violence. *Proceedings of the National Academy of Sciences*. `https://papers.ssrn.com/sol3/papers.cfm?abstract_id=3431132`.

Knox, Dean , Teppei Yamamoto, Matthew A. Baum, and Adam J. Berinsky (2019). Design, identification, and sensitivity analysis for patient preference trials. *Journal of the American Statistical Association 00*(0), 1–15.

Kocieniewski, David (2002, March 21). Study suggests racial gap in speeding in new jersey. *The New York Times*. `https://www.nytimes.com/2002/03/21/nyregion/study-suggests-racial-gap-in-speeding-in-new-jersey.html`.

Lange, James E. , Mark B. Johnson, and Robert B. Voas (2005). Testing the racial profiling hypothesis for seemingly disparate traffic stops on the new jersey turnpike. *Justice Quarterly 22*(2), 193–223.

Lasswell, Harold D. (1936). *Politics: Who Gets What, When, How*. Pickle Partners Publishing.

Lee, David S. (2009). Training, wages, and sample selection: Estimating sharp bounds on treatment effects. *The Review of Economic Studies 76*(3), 1071–1102.

Lerman, Amy and Vesla Weaver (2014). *Arresting Citizenship: The Democratic Consequences of American Crime Control*. University of Chicago Press.

Lipsky, Michael (1980). *Street-Level Bureaucracy: Dilemmas of the Individual in Public Service*. Russell Sage Foundation.

Magaloni, Beatriz , Edgar Franco, and Vanessa Melo (2015). Killing in the slums: an impact evaluation of police reform in rio de janeiro. *Working Paper* (556). `https://siepr.stanford.edu/sites/default/files/publications/556wp_9.pdf`.

Manski, C. F. (1995). *Identification Problems in the Social Sciences*. Harvard University Press.

Mummolo, Jonathan (2018a). Militarization fails to enhance police safety or reduce crime but may harm police reputation. *Proceedings of the National Academy of Sciences of the United States of America 115*(37), 9181–9186.

Mummolo, Jonathan (2018b). Modern police tactics, police-citizen interactions and the prospects for reform. *Journal Of Politics 80*(1), 1–15.

Nix, Justin , Bradley A. Campbell, Edward H. Byers, and Geoffrey P. Alpert (2017). A bird's eye view of civilians killed by police in 2015 further evidence of implicit bias. *Criminology & Public Policy 16*(1), 309–340.

Nyhan, Brendan , Christopher Skovron, and Rocío Titiunik (2017). Differential registration bias in voter file data: A sensitivity analysis approach. *American Journal of Political Science 61*(3), 744–760.

NYPD (2017). Use of force report, 2017. Technical report. `https://www1.nyc.gov/assets/nypd/downloads/pdf/use-of-force/use-of-force-2017.pdf`.

Orne, Martin T. (1962). On the social psychology of the psychological experiment: With particular reference to demand characteristics and their implications. *American Psychologist 17*(1), 776–783.

Ostrom, Elinor and Gordon Whitaker (1973). Does local community control of police make a difference? some preliminary findings. *American Journal of Political Science 17*(1), 48–76.

Pearl, Judea (2001). Direct and indirect effects. *Proceedings of the seventeenth conference on uncertainty in artificial intelligence*, 411–420.

Pearl, Judea (2018). Does obesity shorten life? or is it the soda? on non-manipulable causes. *Journal of Causal Inference 6*(2), 1–7.

Peyton, Kyle , Michael Sierra-Arévalo, and David G. Rand (2019). A field experiment on community policing and police legitimacy. *Proceedings of the National Academy of Sciences 116*(40), 19894–19898.

Phelps, Edmund S. (1972). The statistical theory of racism and sexism. *American Economic Review 62*(1), 659–661.

Ridgeway, Greg (2006). Assessing the effect of race bias in post-traffic stop outcomes using propensity scores. *Journal of Quantitative Criminology 22*(1), 1–29.

Ridgeway, Greg and John MacDonald (2010). *Race, Ethnicity, and Policing: New and Essential Readings*, Chapter Methods for Assessing Racially Biased Policing. NYU Press.

Robins, James M. , Miguel A. Hernán, and Babette Brumback (2000). Marginal structural models and causal inference in epidemiology. *Epidemiology 11*(5), 550–60.

Rosenbaum, Paul R. (1984). The consequences of adjustment for a concomitant variable that has been affected by the treatment. *Journal of the Royal Statistical Society 147*(5), 656–666.

Rubin, Donald B. (1974). Estimating causal effects of treatments in randomized and non-randomized studies. *Journal of Educational Psychology 66*(5), 688–701.

Rubin, Donald B. (1990). Formal mode of statistical inference for causal effects. *Journal of statistical planning and inference 25*(3), 279–292.

Rubin, Donald B. (2000). Causal inference without counterfactuals: Comment. *Journal of the American Statistical Association 95*(450), 435–438.

Samii, Cyrus (2016). Causal empiricism in quantitative research. *The Journal of Politics 78*(3), 941–955.

Sen, Maya and Omar Wasow (2016). Race as a bundle of sticks: Designs that estimate effects of seemingly immutable characteristics. *Annual Review of Political Science 19*, 499–522.

Simoiu, Camelia , Sam Corbett-Davies, and Sharad Goel (2017). The problem of infra-marginality in outcome tests for discrimination. *The Annals of Applied Statistics 11*(3), 1193–1216. `https://arxiv.org/abs/1706.05678`.

Smith, Douglas A. , Christy A. Visher, , and Laura A. Davidson (1984). Equity and discretionary justice: The influence of race on police arrest decisions. *J. Crim. L. & Criminology 75*(1).

Soss, Joe and Vesla Weaver (2017). Police are our government: Politics, political science, and the policing of race–class subjugated communities. *Annual Review of Political Science 20*, 565–591.

VanderWeele, Tyler J. (2009). Marginal structural models for the estimation of direct and indirect effects. *Epidemiology 20*(1), 18–26.

VanderWeele, Tyler J. (2011). Principal stratification–uses and limitations. *The International Journal of Biostatistics 7*(1), 1–14.

West, Jeremy (2018). Racial bias in police investigations. Working Paper `https://people.ucsc.edu/~jwest1/articles/West_RacialBiasPolice.pdf`.

White, Ariel (2019). Misdemeanor disenfranchisement? the demobilizing effects of brief jail spells on potential voters. *American Political Science Review 113*(2), 311–324.

Wilson, James Q. (1968). *Varieties of Police Behavior*. Harvard University Press.

Wilson, James Q (1989). *Bureaucracy: What Government Agencies Do and Why They Do It*. Basic Books.

Yamamoto, Teppei (2012). Understanding the past: Statistical analysis of causal attribution. *American Journal of Political Science 56*(1), 237–256.

Zhang, Junni L. and Donald B. Rubin (2003). Estimation of causal effects via principal stratification when some outcomes are truncated by "death". *Journal of Educational and Behavioral Statistics 28*(4), 353–368.

Administrative Records Mask Racially Biased Policing

Online Appendix

Contents

A   Detailed proofs

A.1   Bias for ATE$_{M=1}$

We first derive the bias of the local difference in means (that is, among encounters with $X_i = x$, $\hat{\Delta}_x = \overline{Y_i|D_i = 1, M_i = 1, X_i = x} - \overline{Y_i|D_i = 0, M_i = 1, X_i = x}$), in estimating the local average treatment effect among stops, ATE$_{M=1,x} = \mathbb{E}[Y_i(1, M_i(1))|M_i = 1, X_i = x] - \mathbb{E}[Y_i(0, M_i(0))|M_i = 1, X_i = x]$. The overall bias is then given by $\sum_x \left(\mathbb{E}[\hat{\Delta}_x] - \text{ATE}_{M=1,x} \right) \Pr(X_i = x|M_i = 1)$.

$\mathbb{E}[\hat{\Delta}_x] - \text{ATE}_{M=1,x}$

$= (\mathbb{E}[\overline{Y_i|D_i = 1, M_i = 1, X_i = x} - \overline{Y_i|D_i = 0, M_i = 1, X_i = x}])$

$\quad - (\mathbb{E}[Y_i(1, M_i(1))|M_i = 1, X_i = x] - \mathbb{E}[Y_i(0, M_i(0))|M_i = 1, X_i = x])$

$= \mathbb{E}[Y_i|D_i = 1, M_i(D_i) = 1, X_i = x] - \mathbb{E}[Y_i|D_i = 0, M_i(D_i) = 1, X_i = x]$

$\quad - \mathbb{E}[Y_i(1, M_i(1))|M_i(D_i) = 1, X_i = x] + \mathbb{E}[Y_i(0, M_i(0))|M_i(D_i) = 1, X_i = x]$

$= \mathbb{E}[Y_i|D_i = 1, M_i(D_i) = 1, X_i = x] - \mathbb{E}[Y_i(1, M_i(1))|M_i(D_i) = 1, X_i = x]$

$\quad - \mathbb{E}[Y_i|D_i = 0, M_i(D_i) = 1, X_i = x] + \mathbb{E}[Y_i(0, M_i(0))|M_i(D_i) = 1, X_i = x]$

$= \mathbb{E}[Y_i(1, M_i(1))|D_i = 1, M_i(D_i) = 1, X_i = x]$

$\quad - \mathbb{E}[Y_i(1, M_i(1))|D_i = 1, M_i(D_i) = 1, X_i = x]\Pr(D_i = 1|M_i(D_i) = 1, X_i = x)$

$\quad - \mathbb{E}[Y_i(1, M_i(1))|D_i = 0, M_i(D_i) = 1, X_i = x]\Pr(D_i = 0|M_i(D_i) = 1, X_i = x)$

$\quad - \mathbb{E}[Y_i(0, M_i(0))|D_i = 0, M_i(D_i) = 1, X_i = x]$

$\quad + \mathbb{E}[Y_i(0, M_i(0))|D_i = 1, M_i(D_i) = 1, X_i = x]\Pr(D_i = 1|M_i(D_i) = 1, X_i = x)$

$\quad + \mathbb{E}[Y_i(0, M_i(0))|D_i = 0, M_i(D_i) = 1, X_i = x]\Pr(D_i = 0|M_i(D_i) = 1, X_i = x)$

$= \mathbb{E}[Y_i(1, M_i(1))|D_i = 1, M_i(D_i) = 1, X_i = x]\Pr(D_i = 0|M_i(D_i) = 1, X_i = x)$

$\quad - \mathbb{E}[Y_i(1, M_i(1))|D_i = 0, M_i(D_i) = 1, X_i = x]\Pr(D_i = 0|M_i(D_i) = 1, X_i = x)$

$\quad - \mathbb{E}[Y_i(0, M_i(0))|D_i = 0, M_i(D_i) = 1, X_i = x]\Pr(D_i = 1|M_i(D_i) = 1, X_i = x)$

$\quad + \mathbb{E}[Y_i(0, M_i(0))|D_i = 1, M_i(D_i) = 1, X_i = x]\Pr(D_i = 1|M_i(D_i) = 1, X_i = x)$

under mediator monotonicity, $\Pr(M_i(1) = 0|D_i = 0, M_i(D_i) = 1, X_i = x) = 0$ and $\Pr(M_i(1) = 1|D_i = 0, M_i(D_i) = 1, X_i = x) = 1$,

$= \mathbb{E}[Y_i(1, M_i(1))|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x]$

$\qquad \Pr(M_i(0) = 1|D_i = 1, M_i(D_i) = 1, X_i = x)\Pr(D_i = 0|M_i(D_i) = 1, X_i = x)$

$\quad + \mathbb{E}[Y_i(1, M_i(1))|D_i = 1, M_i(1) = 1, M_i(0) = 0, X_i = x]$

$$\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)\Pr(D_i = 0|M_i(D_i) = 1, X_i = x)$$

$$- \mathbb{E}[Y_i(1, M_i(1))|D_i = 0, M_i(1) = 1, M_i(0) = 1, X_i = x]$$

~~$$\Pr(M_i(1) = 1|D_i = 0, M_i(D_i) = 1, X_i = x)\Pr(D_i = 0|M_i(D_i) = 1, X_i = x)$$~~

~~$$- \mathbb{E}[Y_i(1, M_i(1))|D_i = 0, M_i(1) = 0, M_i(0) = 1, X_i = x]$$~~

~~$$\Pr(M_i(1) = 0|D_i = 0, M_i(D_i) = 1, X_i = x)\Pr(D_i = 0|M_i(D_i) = 1, X_i = x)$$~~

$$- \mathbb{E}[Y_i(0, M_i(0))|D_i = 0, M_i(1) = 1, M_i(0) = 1, X_i = x]$$

~~$$\Pr(M_i(1) = 1|D_i = 0, M_i(D_i) = 1, X_i = x)\Pr(D_i = 1|M_i(D_i) = 1, X_i = x)$$~~

~~$$- \mathbb{E}[Y_i(0, M_i(0))|D_i = 0, M_i(1) = 0, M_i(0) = 1, X_i = x]$$~~

~~$$\Pr(M_i(1) = 0|D_i = 0, M_i(D_i) = 1, X_i = x)\Pr(D_i = 1|M_i(D_i) = 1, X_i = x)$$~~

$$+ \mathbb{E}[Y_i(0, M_i(0))|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x]$$

$$\Pr(M_i(0) = 1|D_i = 1, M_i(D_i) = 1, X_i = x)\Pr(D_i = 1|M_i(D_i) = 1, X_i = x)$$

$$+ \mathbb{E}[Y_i(0, M_i(0))|D_i = 1, M_i(1) = 1, M_i(0) = 0, X_i = x]$$

$$\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)\Pr(D_i = 1|M_i(D_i) = 1, X_i = x)$$

$$= \mathbb{E}[Y_i(1, M_i(1))|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x]$$

$$\Pr(M_i(0) = 1|D_i = 1, M_i(D_i) = 1, X_i = x)$$

$$- \mathbb{E}[Y_i(1, M_i(1))|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x]$$

$$\Pr(M_i(0) = 1|D_i = 1, M_i(D_i) = 1, X_i = x)\Pr(D_i = 1|M_i(D_i) = 1, X_i = x)$$

$$+ \mathbb{E}[Y_i(1, M_i(1))|D_i = 1, M_i(1) = 1, M_i(0) = 0, X_i = x]$$

$$\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)$$

$$- \mathbb{E}[Y_i(1, M_i(1))|D_i = 1, M_i(1) = 1, M_i(0) = 0, X_i = x]$$

$$\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)\Pr(D_i = 1|M_i(D_i) = 1, X_i = x)$$

$$- \mathbb{E}[Y_i(1, M_i(1))|D_i = 0, M_i(1) = 1, M_i(0) = 1, X_i = x]$$

$$+ \mathbb{E}[Y_i(1, M_i(1))|D_i = 0, M_i(1) = 1, M_i(0) = 1, X_i = x]$$

$$\Pr(D_i = 1|M_i(D_i) = 1, X_i = x)$$

$$- \mathbb{E}[Y_i(0, M_i(0))|D_i = 0, M_i(1) = 1, M_i(0) = 1, X_i = x]$$

$$\Pr(D_i = 1|M_i(D_i) = 1, X_i = x)$$

$$+ \mathbb{E}[Y_i(0, M_i(0))|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x]$$

$$\Pr(M_i(0) = 1|D_i = 1, M_i(D_i) = 1, X_i = x)\Pr(D_i = 1|M_i(D_i) = 1, X_i = x)$$

$$+ \mathbb{E}[Y_i(0, M_i(0))|D_i = 1, M_i(1) = 1, M_i(0) = 0, X_i = x]$$

$$\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)\Pr(D_i = 1|M_i(D_i) = 1, X_i = x)$$

adding and subtracting $\mathbb{E}[Y_i(1, M_i(1))|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x]\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)$,

$$= \mathbb{E}[Y_i(1, M_i(1)) - Y_i(0, M_i(0))|D_i = 0, M_i(1) = 1, M_i(0) = 1, X_i = x]\Pr(D_i = 1|M_i(D_i) = 1, X_i = x)$$

$$- \mathbb{E}[Y_i(1, M_i(1)) - Y_i(0, M_i(0))|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x]$$

$$\Pr(M_i(0) = 1|D_i = 1, M_i(D_i) = 1, X_i = x)\Pr(D_i = 1|M_i(D_i) = 1, X_i = x)$$

$$- \mathbb{E}[Y_i(1, M_i(1)) - Y_i(0, M_i(0))|D_i = 1, M_i(1) = 1, M_i(0) = 0, X_i = x]$$

$$\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)\Pr(D_i = 1|M_i(D_i) = 1, X_i = x)$$

$$- \mathbb{E}[Y_i(1, M_i(1))|D_i = 0, M_i(1) = 1, M_i(0) = 1, X_i = x]$$

$$+ \mathbb{E}[Y_i(1, M_i(1))|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x]\Pr(M_i(0) = 1|D_i = 1, M_i(D_i) = 1, X_i = x)$$

$$+ \mathbb{E}[Y_i(1, M_i(1))|D_i = 1, M_i(1) = 1, M_i(0) = 0, X_i = x]\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)$$

$$+ \mathbb{E}[Y_i(1, M_i(1))|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x]\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)$$

$$- \mathbb{E}[Y_i(1, M_i(1))|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x]\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)$$

substituting potential mediators based on principal strata,

$$= \mathbb{E}[Y_i(1, 1) - Y_i(0, 1)|D_i = 0, M_i(1) = 1, M_i(0) = 1, X_i = x]\Pr(D_i = 1|M_i(D_i) = 1, X_i = x)$$

$$- \mathbb{E}[Y_i(1, 1) - Y_i(0, 1)|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x]$$

$$\Pr(M_i(0) = 1|D_i = 1, M_i(D_i) = 1, X_i = x)\Pr(D_i = 1|M_i(D_i) = 1, X_i = x)$$

$$- \mathbb{E}[Y_i(1, 1) - Y_i(0, 0)|D_i = 1, M_i(1) = 1, M_i(0) = 0, X_i = x]$$

$$\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)\Pr(D_i = 1|M_i(D_i) = 1, X_i = x)$$

$$+ \mathbb{E}[Y_i(1, 1)|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x]$$

$$- \mathbb{E}[Y_i(1, 1)|D_i = 0, M_i(1) = 1, M_i(0) = 1, X_i = x]$$

$$+ \mathbb{E}[Y_i(1, 1)|D_i = 1, M_i(1) = 1, M_i(0) = 0, X_i = x]\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)$$

$$- \mathbb{E}[Y_i(1, 1)|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x]\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)$$

under mandatory reporting, $Y_i(d, 0) = 0$,

$$= \mathbb{E}[Y_i(1, 1) - Y_i(0, 1)|D_i = 0, M_i(1) = 1, M_i(0) = 1, X_i = x]\Pr(D_i = 1|M_i(D_i) = 1, X_i = x)$$

$$- \mathbb{E}[Y_i(1, 1) - Y_i(0, 1)|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x]$$

$$\Pr(M_i(0) = 1|D_i = 1, M_i(D_i) = 1, X_i = x)\Pr(D_i = 1|M_i(D_i) = 1, X_i = x)$$

$$- \mathbb{E}[Y_i(1, 1)|D_i = 1, M_i(1) = 1, M_i(0) = 0, X_i = x]$$

$$\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)\Pr(D_i = 1|M_i(D_i) = 1, X_i = x)$$

$$+ \mathbb{E}[Y_i(1,1)|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x]$$

$$- \mathbb{E}[Y_i(1,1)|D_i = 0, M_i(1) = 1, M_i(0) = 1, X_i = x]$$

$$+ \mathbb{E}[Y_i(1,1)|D_i = 1, M_i(1) = 1, M_i(0) = 0, X_i = x]\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)$$

$$- \mathbb{E}[Y_i(1,1)|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x]\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)$$

invoking assumption 4(b) (treatment ignorability),

$$= \Big(\mathbb{E}[Y_i(1,1) - Y_i(0,1)|M_i(1) = 1, M_i(0) = 1, X_i = x]$$

$$- \mathbb{E}[Y_i(1,1) - Y(0,0)|M_i(1) = 1, M_i(0) = 0, X_i = x]$$

$$\Big) \Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)\Pr(D_i = 1|M_i(D_i) = 1, X_i = x)$$

$$- (\mathbb{E}[Y_i(1,1)|M_i(1) = 1, M_i(0) = 1, X_i = x]$$

$$- \mathbb{E}[Y_i(1,1)|M_i(1) = 1, M_i(0) = 0, X_i = x])\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)$$

which is Equation 6.

A.2   Bias for ATT$_{M=1}$

Next, we consider the bias that results when the local difference in means is used as an estimator for the local average racial effect among stopped minorities, ATT$_{M=1,x}$ = $\mathbb{E}[Y_i(1, M_i(1))|D_i = 1, M_i = 1, X_i = x] - \mathbb{E}[Y_i(0, M_i(0))|D_i = 1, M_i = 1, X_i = x]$. Again, overall bias is found by the weighted average of local biases, $\sum_x \Big(\mathbb{E}[\hat{\Delta}_x] - \text{ATT}_{M=1,x} \Big) \Pr(X_i = x|D_i = 1, M_i = 1)$.

$$\mathbb{E}[\hat{\Delta}_x] - \text{ATT}_{M=1,x} = (\mathbb{E}[\overline{Y_i|D_i = 1, M_i = 1, X_i = x} - \overline{Y_i|D_i = 0, M_i = 1, X_i = x}])$$

$$- (\mathbb{E}[Y_i(1, M_i(1))|D_i = 1, M_i = 1, X_i = x] - \mathbb{E}[Y_i(0, M_i(0))|D_i = 1, M_i = 1, X_i = x])$$

$$= \mathbb{E}[Y_i(1, M_i(1))|D_i = 1, M_i(D_i) = 1, X_i = x]$$

$$- \mathbb{E}[Y_i(0, M_i(0))|D_i = 0, M_i(D_i) = 1, X_i = x]$$

$$- \mathbb{E}[Y_i(1, M_i(1))|D_i = 1, M_i(D_i) = 1, X_i = x]$$

$$+ \mathbb{E}[Y_i(0, M_i(0))|D_i = 1, M_i(D_i) = 1, X_i = x]$$

$$= - \mathbb{E}[Y_i(0, M_i(0))|D_i = 0, M_i(D_i) = 1, X_i = x]$$

$$+ \mathbb{E}[Y_i(0, M_i(0))|D_i = 1, M_i(D_i) = 1, X_i = x]$$

$$= - \mathbb{E}[Y_i(0,1)|D_i = 0, M_i(1) = 1, M_i(0) = 1, X_i = x]\Pr(M_i(1) = 1|D_i = 0, M_i(D_i) = 1, X_i = x)$$

$$- \mathbb{E}[Y_i(0,1)|D_i = 0, M_i(1) = 0, M_i(0) = 1, X_i = x]\Pr(M_i(1) = 0|D_i = 0, M_i(D_i) = 1, X_i = x)$$

$$+ \mathbb{E}[Y_i(0,1)|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x]\Pr(M_i(0) = 1|D_i = 1, M_i(D_i) = 1, X_i = x)$$

$$+ \mathbb{E}[Y_i(0,0)|D_i = 1, M_i(1) = 1, M_i(0) = 0, X_i = x]\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)$$

by mediator monotonicity

$$
\begin{aligned}
= \ & -\mathbb{E}[Y_i(0,1)|D_i = 0, M_i(1) = 1, M_i(0) = 1, X_i = x] \\
& + \mathbb{E}[Y_i(0,1)|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x]\Pr(M_i(0) = 1|D_i = 1, M_i(D_i) = 1, X_i = x) \\
& + \mathbb{E}[Y_i(0,0)|D_i = 1, M_i(1) = 1, M_i(0) = 0, X_i = x]\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)
\end{aligned}
$$

by mandatory reporting

$$
\begin{aligned}
= \ & -\mathbb{E}[Y_i(0,1)|D_i = 0, M_i(1) = 1, M_i(0) = 1, X_i = x] \\
& + \mathbb{E}[Y_i(0,1)|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x]\Pr(M_i(0) = 1|D_i = 1, M_i(D_i) = 1, X_i = x)
\end{aligned}
$$

by treatment ignorability

$$
= -\mathbb{E}[Y_i(0,1)|M_i(1) = 1, M_i(0) = 1, X_i = x]\Pr(M_i(0) = 0|M_i(1) = 1, X_i = x)
$$

A.3   Bias for CDE$_{M=1}$

The CDE$_{M=1}$ is defined as

$$
\text{CDE}_{M=1} = \mathbb{E}[Y_i(1,1)|M_i(D_i) = 1] - \mathbb{E}[Y_i(0,1)|M_i(D_i) = 1] \tag{1}
$$

It is a somewhat contrived estimand because it necessarily involves counterfactuals that, for racial-stop encounters, could never realize even if researchers could somehow randomize civilian race in police-civilian encounters. For example, when a minority civilian is racially stopped for a "furtive movement" and reaches for their wallet, it makes little sense to consider an officer's potential use of force if the civilian suddenly became white at that moment: had police observed a white civilian from the onset, a stop would never have occurred. Moreover, the assumptions required for such counterfactuals are fundamentally unverifiable (Robins and Greenland, 1992), unless the experimentalist can somehow also manipulate officer stopping decisions without distorting outcomes. (We note that always-stop encounters are not subject to this issue, but analysts cannot hope to estimate the conditional CDE in this group because it is impossible to identify which minority encounters belong to this group.)

In the previous bias expression for the ATE$_{M=1}$, white individuals in the data—necessarily belonging to the always-stop group, $M_i(1) = M_i(0) = 1$—were used to estimate the $Y_i(0, M_i(0))$ potential outcomes of minority encounters in the data. Unavoidable bias arose as long as any minority individuals in the data belonged to the racial-stop group—had these individuals been white,

they would never have been stopped, and hence would not be subject to force. Changing the target estimand to the $CDE_{M=1}$ conceptually sidesteps this specific issue by considering a different counterfactual, $Y_i(0,1)$ instead of $Y_i(0, M_i(0))$. But as we note above, for encounters in which only minority civilians would be stopped—that is, encounters with $M_i(1) = 1$ and $M_i(0) = 0$—this new counterfactual represents an impossible cross-world scenario. The $CDE_{M=1}$ asks whether force would have been used if officers were forced, against nature, to stop a white individual in this encounter as if they were a minority.

We now demonstrate that the local difference in means remains biased for the local controlled direct effect, $CDE_{M=1,x} = \mathbb{E}[Y_i(1,1)|M_i = 1, X_i = x] - \mathbb{E}[Y_i(0,1)|M_i = 1, X_i = x]$, unless officers are as violent toward minorities in always-stop encounters (where they are forced to intervene) as they are in racially discriminatory stops (where they are free to exercise discretion). In other words, for the naïve estimator to recover the $CDE_{M=1}$, Assumption 5 must hold. The derivation is almost identical to that of the $ATE_{M=1,x}$, differing only in that all individuals are held at $M_i = 1$ instead of allowed stops to vary with civilian race, $M_i(D_i)$. Bias for $CDE_{M=1}$ is then given by the weighted average of local biases, $\sum_x \left(\mathbb{E}[\hat{\Delta}_x] - CDE_{M=1,x} \right) \Pr(X_i = x|M_i = 1)$.

$$
\begin{aligned}
\mathbb{E}[\hat{\Delta}_x] - CDE_{M=1,x} &= (\mathbb{E}[\overline{Y_i|D_i = 1, M_i = 1, X_i = x} - \overline{Y_i|D_i = 0, M_i = 1, X_i = x}]) \\
&\quad - (\mathbb{E}[Y_i(1,1)|M_i = 1, X_i = x] - \mathbb{E}[Y_i(0,1)|M_i = 1, X_i = x]) \\
&= \mathbb{E}[Y_i|D_i = 1, M_i(D_i) = 1, X_i = x] - \mathbb{E}[Y_i|D_i = 0, M_i(D_i) = 1, X_i = x] \\
&\quad - \mathbb{E}[Y_i(1,1)|M_i(D_i) = 1, X_i = x] + \mathbb{E}[Y_i(0,1)|M_i(D_i) = 1, X_i = x] \\
&= \mathbb{E}[Y_i(1,1) - Y_i(0,1)|D_i = 0, M_i(1) = 1, M_i(0) = 1, X_i = x]\Pr(D_i = 1|M_i(D_i) = 1, X_i = x) \\
&\quad - \mathbb{E}[Y_i(1,1) - Y_i(0,1)|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x] \\
&\qquad \Pr(M_i(0) = 1|D_i = 1, M_i(D_i) = 1, X_i = x)\Pr(D_i = 1|M_i(D_i) = 1, X_i = x) \\
&\quad - \mathbb{E}[Y_i(1,1) - Y_i(0,1)|D_i = 1, M_i(1) = 1, M_i(0) = 0, X_i = x] \\
&\qquad \Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)\Pr(D_i = 1|M_i(D_i) = 1, X_i = x) \\
&\quad + \mathbb{E}[Y_i(1,1)|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x] \\
&\quad - \mathbb{E}[Y_i(1,1)|D_i = 0, M_i(1) = 1, M_i(0) = 1, X_i = x] \\
&\quad + \mathbb{E}[Y_i(1,1)|D_i = 1, M_i(1) = 1, M_i(0) = 0, X_i = x]\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x) \\
&\quad - \mathbb{E}[Y_i(1,1)|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x]\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)
\end{aligned}
$$

under assumption 4(b),

$$
\begin{aligned}
&= \Big(\mathbb{E}[Y_i(1,1) - Y_i(0,1)|M_i(1) = 1, M_i(0) = 1, X_i = x] \\
&\quad - \mathbb{E}[Y_i(1,1) - Y_i(0,1)|M_i(1) = 1, M_i(0) = 0, X_i = x]
\end{aligned}
$$

$$) \Pr(M_i(0) = 0 | D_i = 1, M_i(D_i) = 1, X_i = x) \Pr(D_i = 1 | M_i(D_i) = 1, X_i = x)$$

$$- (\mathbb{E}[Y_i(1,1) | M_i(1) = 1, M_i(0) = 1, X_i = x] - \mathbb{E}[Y_i(1,1) | M_i(1) = 1, M_i(0) = 0, X_i = x])$$

$$\Pr(M_i(0) = 0 | D_i = 1, M_i(D_i) = 1, X_i = x)$$

which reproduces Equation 7.

Finally, we demonstrate that this bias is weakly negative. Rearranging terms yields

$$\mathbb{E}[Y_i(1,1) - Y_i(0,1) | M_i(1) = 1, M_i(0) = 1, X_i = x]$$

$$\times \Pr(M_i(0) = 0 | D_i = 1, M_i(D_i) = 1, X_i = x) \Pr(D_i = 1 | M_i(D_i) = 1, X_i = x)$$

$$- \mathbb{E}[Y_i(1,1) - Y_i(0,1) | M_i(1) = 1, M_i(0) = 0, X_i = x]$$

$$\times \Pr(M_i(0) = 0 | D_i = 1, M_i(D_i) = 1, X_i = x) \Pr(D_i = 1 | M_i(D_i) = 1, X_i = x)$$

$$- \mathbb{E}[Y_i(1,1) | M_i(1) = 1, M_i(0) = 1, X_i = x]$$

$$\times \Pr(M_i(0) = 0 | D_i = 1, M_i(D_i) = 1, X_i = x)$$

$$+ \mathbb{E}[Y_i(1,1) | M_i(1) = 1, M_i(0) = 0, X_i = x]$$

$$\times \Pr(M_i(0) = 0 | D_i = 1, M_i(D_i) = 1, X_i = x)$$

$$= -\mathbb{E}[Y_i(0,1) | M_i(1) = 1, M_i(0) = 1, X_i = x]$$

$$\times \Pr(M_i(0) = 0 | D_i = 1, M_i(D_i) = 1, X_i = x) \Pr(D_i = 1 | M_i(D_i) = 1, X_i = x)$$

$$+ \mathbb{E}[Y_i(0,1) | M_i(1) = 1, M_i(0) = 0, X_i = x]$$

$$\times \Pr(M_i(0) = 0 | D_i = 1, M_i(D_i) = 1, X_i = x) \Pr(D_i = 1 | M_i(D_i) = 1, X_i = x)$$

$$- \mathbb{E}[Y_i(1,1) | M_i(1) = 1, M_i(0) = 1, X_i = x]$$

$$\times \Pr(M_i(0) = 0 | D_i = 1, M_i(D_i) = 1, X_i = x)(1 - \Pr(D_i = 1 | M_i(D_i) = 1, X_i = x))$$

$$+ \mathbb{E}[Y_i(1,1) | M_i(1) = 1, M_i(0) = 0, X_i = x]$$

$$\times \Pr(M_i(0) = 0 | D_i = 1, M_i(D_i) = 1, X_i = x)(1 - \Pr(D_i = 1 | M_i(D_i) = 1, X_i = x))$$

The sum of the first and second terms is weakly negative by Assumption 3, because the magnitude of the first term is greater than that of the second, and similarly the sum of the third and fourth terms is also weakly negative. Therefore, bias is weakly negative when the naïve estimator is used to estimate the $\mathrm{CDE}_{M=1}$.

A.4 Nonparametric sharp bounds for $\mathbf{ATE}_{M=1}$

In this section, we derive nonparametric sharp bounds for the $\mathrm{ATE}_{M=1,x}$. We begin with the case when the proportion of racially discriminatory stops among reported minority encounters, $\Pr(M_i(0) = 0 | D_i = 1, M_i(D_i) = 1, X_i = x)$, is known or can be assumed. Rearrangement of Equa-

tion 6 (within levels of X) yields

$$\text{ATE}_{M=1,x} = \mathbb{E}[\hat{\Delta}_x]$$

$$+ \mathbb{E}[Y_i(1,1)|M_i(1) = 1, M_i(0) = 1, X_i = x]$$

$$\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)(1 - \Pr(D_i = 1|M_i(D_i) = 1, X_i = x))$$

$$- \mathbb{E}[Y_i(1,1)|M_i(1) = 1, M_i(0) = 0, X_i = x]$$

$$\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)(1 - \Pr(D_i = 1|M_i(D_i) = 1, X_i = x))$$

$$+ \mathbb{E}[Y_i(0,1)|M_i(1) = 1, M_i(0) = 1, X_i = x]$$

$$\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)\Pr(D_i = 1|M_i(D_i) = 1, X_i = x)$$

$$= \mathbb{E}[\hat{\Delta}_x]$$

$$+ \frac{\mathbb{E}[Y_i(1,1)|D_i = 1, M_i(1) = 1, X_i = x]}{\Pr(M_i(0) = 1|D_i = 1, M_i(1) = 1, X_i = x)}\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)\Pr(D_i = 0|M_i(D_i) = 1, X_i = x)$$

$$- \frac{\mathbb{E}[Y_i(1,1)|D_i = 1, M_i(1) = 1, M_i(0) = 0, X_i = x]}{\Pr(M_i(0) = 1|D_i = 1, M_i(1) = 1, X_i = x)}\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)^2 \qquad (2)$$

$$\Pr(D_i = 0|M_i(D_i) = 1, X_i = x)$$

$$- \mathbb{E}[Y_i(1,1)|M_i(1) = 1, M_i(0) = 0, X_i = x]\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)\Pr(D_i = 0|M_i(D_i) = 1, X_i = x)$$

$$+ \mathbb{E}[Y_i(0,1)|M_i(1) = 1, M_i(0) = 1, X_i = x]\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)\Pr(D_i = 1|M_i(D_i) = 1, X_i = x)$$

$$= \mathbb{E}[\hat{\Delta}_x]$$

$$+ \frac{\mathbb{E}[Y_i|D_i = 1, M_i = 1, X_i = x]}{\Pr(M_i(0) = 1|D_i = 1, M_i(1) = 1, X_i = x)}\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)\Pr(D_i = 0|M_i = 1, X_i = x)$$

$$- \frac{\mathbb{E}[Y_i(1,1)|D_i = 1, M_i(1) = 1, M_i(0) = 0, X_i = x]}{\Pr(M_i(0) = 1|D_i = 1, M_i(1) = 1, X_i = x)}\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)^2 \qquad (3)$$

$$\Pr(D_i = 0|M_i = 1, X_i = x)$$

$$- \mathbb{E}[Y_i(1,1)|M_i(1) = 1, M_i(0) = 0, X_i = x]\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)\Pr(D_i = 0|M_i = 1, X_i = x)$$

$$+ \mathbb{E}[Y_i|D_i = 0, M_i = 1, X_i = x]\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)\Pr(D_i = 1|M_i = 1, X_i = x)$$

$$(4)$$

We then construct bounds on $\mathbb{E}[Y_i(1,1)|D_i = 1, M_i(1) = 1, M_i(0) = 0, X_i = x]$ based on Fréchet inequalities for the joint distribution, $\Pr(Y_i(1,1) = 1, M_i(0) = 0|D_i = 1, M_i(1) = 1, X_i = x)$, which incorporate marginal information about $Y_i(1,1)$ and $M_i(0)$.

$$\frac{\max\{0, \Pr(M_i(0) = 0|D_i = 1, M_i(1) = 1, X_i = x) + \mathbb{E}[Y_i|D_i = 1, M_i = 1, X_i = x] - 1\}}{\Pr(M_i(0) = 0|D_i = 1, M_i(1) = 1, X_i = x)}$$

$$\leq \quad \mathbb{E}[Y_i(1,1)|D_i = 1, M_i(1) = 1, M_i(0) = 0, X_i = x] \quad \leq \qquad (5)$$

$$\frac{\min\{\Pr(M_i(0) = 0|D_i = 1, M_i(1) = 1, X_i = x), \mathbb{E}[Y_i|D_i = 1, M_i = 1, X_i = x]\}}{\Pr(M_i(0) = 0|D_i = 1, M_i(1) = 1, X_i = x)}$$

These bounds are sharp given only marginal information, $\Pr(Y_i(1,1) = 1|D_i = 1, M_i(1) = 1, X_i = x)$

and $\Pr(M_i(0) = 0|D_i = 1, M_i(1) = 1, X_i = x)$. However, the upper bound can be tightened further under Assumption 3, which implies $\mathbb{E}[Y_i(1,1)|D_i = 1, M_i(1) = 1, M_i(0) = 0, X_i = x] \leq \mathbb{E}[Y_i|D_i = 1, M_i = 1, X_i = x]$; this is at least as small as the upper Fréchet bound.

Finally, note that the reported data contain no information that can be used to constrain the proportion of racially discriminatory minority stops, $\Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)$. If this proportion were zero, then the distribution of civilian race in police reports would reflect that of all police encounters (within levels of X). The reported data cannot distinguish between this possibility and an alternative population in which $\rho_x = \Pr(M_i(0) = 0|D_i = 1, M_i(D_i) = 1, X_i = x)$ is large, but white encounters are also larger by the proportion $1/(1 - \rho_x)$. Without side information about the total number of encounters, this proportion can take on any value in $[0, 1)$. Therefore, sharp bounds on $\text{ATE}_{M=1}$ alone are obtained by substituting Equation 5 into Equation 4 and setting the proportion of racial stops to unity. The bivariate bounds define the region in which $(\text{ATE}_{M=1}, \rho_x)$ pairs are consistent with the observed data. When ρ_x is set to zero or one, these respectively recover the difference in reported means and the marginal upper bounds on $\text{ATE}_{M=1}$. For $\rho_x \in (0, 1)$,

$$\mathbb{E}[\hat{\Delta}_x] + \rho_x \, \mathbb{E}[Y_i|D_i = 0, M_i = 1, X_i = x](1 - \Pr(D_i = 0|M_i = 1, X_i = x))$$

$$\leq \ \text{ATE}_{M=1,x} \ \leq$$

$$\mathbb{E}[\hat{\Delta}_x]$$

$$+ \frac{\rho_x}{1 - \rho_x}\left(\mathbb{E}[Y_i|D_i = 1, M_i = 1, X_i = x] - \max\left\{0, 1 + \frac{1}{\rho_x}\mathbb{E}[Y_i|D_i = 1, M_i = 1, X_i = x] - \frac{1}{\rho_x}\right\}\right)\Pr(D_i = 0|M_i = 1, X_i = x)$$

$$+ \rho_x \, \mathbb{E}[Y_i|D_i = 0, M_i = 1, X_i = x]\,(1 - \Pr(D_i = 0|M_i = 1, X_i = x)),$$

which reduces to Proposition 1 in the no-covariate case. Otherwise, bounds on $\text{ATE}_{M=1}$ are given by $\sum_x \underline{\text{ATE}}_{M=1,x} \Pr(X_i = x|M_i = 1) \leq \text{ATE}_{M=1} \leq \sum_x \overline{\text{ATE}}_{M=1,x} \Pr(X_i = x|M_i = 1)$, where $\underline{\text{ATE}}_{M=1,x}$ $(\overline{\text{ATE}}_{M=1,x})$ denote the lower (upper) bounds on the local average treatment effect.

Finally, we note that per Equation 3, the $\text{ATT}_{M=1,x}$ can be written

$$\text{ATT}_{M=1,x} = \mathbb{E}[Y_i(1, M_i(1)) - Y_i(0, M_i(0))|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x]\Pr(M_i(0) = 1|D_i = 1, M_i = 1, X_i = x)$$

$$+ \mathbb{E}[Y_i(1, M_i(1)) - Y_i(0, M_i(0))|D_i = 1, M_i(1) = 1, M_i(0) = 0, X_i = x]\Pr(M_i(0) = 0|D_i = 1, M_i = 1, X_i = x)$$

$$= \mathbb{E}[Y_i(1, 1)|D_i = 1, M_i = 1, X_i = x]$$

$$- \mathbb{E}[Y_i(0, 1)|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x]\Pr(M_i(0) = 1|D_i = 1, M_i = 1, X_i = x)$$

$$- \mathbb{E}[Y_i(0, 0)|D_i = 1, M_i(1) = 1, M_i(0) = 0, X_i = x]\Pr(M_i(0) = 0|D_i = 1, M_i = 1, X_i = x)$$

under Assumption 1,

$$= \mathbb{E}[Y_i(1, 1) - |D_i = 1, M_i = 1, X_i = x]$$

$$- \mathbb{E}[Y_i(0, 1)|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x]\Pr(M_i(0) = 1|D_i = 1, M_i = 1, X_i = x)$$

9

and under Assumption 4,

$$= \mathbb{E}[Y_i | D_i = 1, M_i = 1, X_i = x] - \mathbb{E}[Y_i | D_i = 0, M_i = 1, X_i = x] \, (1 - \Pr(M_i(0) = 0 | D_i = 1, M_i = 1, X_i = x))$$

which can be estimated from observed data if the proportion of racial stops is known. It then follows that

$$\mathrm{ATT}_{M=1} = \sum_x \left(\mathbb{E}[Y_i | D_i = 1, M_i = 1, X_i = x] - \mathbb{E}[Y_i | D_i = 0, M_i = 1, X_i = x] + \rho_x \mathbb{E}[Y_i | D_i = 0, M_i = 1, X_i = x] \right)$$

$$= \sum_x \left(\mathbb{E}[\Delta_x] + \rho_x \mathbb{E}[Y_i | D_i = 0, M_i = 1, X_i = x] \right).$$

A.5 Uncertainty of bounds

Here, we describe our approach for constructing confidence intervals for the bounds on these causal quantities. We take X_i, D_i and M_i as fixed, so that uncertainty in the bounds arises strictly from the estimation of the conditional expectations, $\mathbb{E}[Y_i | D_i = d, M_i = 1, X_i = x]$. The asymptotic distribution of the estimated lower and upper bounds endpoints, $(\hat{\underline{\mathrm{ATE}}}_{M=1}, \hat{\overline{\mathrm{ATE}}}_{M=1})$, then follows directly from the asymptotic joint distribution of $\hat{\mathbb{E}}[Y_i | D_i = d, M_i = 1, X_i = x]$ for all d and x. We approximate this through a Monte Carlo simulation in which parameters of the logistic regression models described in Section 5 are sampled from a multivariate normal distribution centered on the parameter estimates and with the estimated covariance matrix. For each parameter sample θ^*, the corresponding bounds endpoint pair $(\underline{\mathrm{ATE}}^*_{M=1}, \overline{\mathrm{ATE}}^*_{M=1})$ is computed deterministically; after drawing a sufficient number of such samples, we numerically obtain the shortest range that fully contains 95% of all simulated bounds intervals. Closely related alternatives to this approach are the bootstrap-based method of Horowitz and Manski (2000) and the fully Bayesian approach taken in Knox et al. (2019). For the analysis in Section 5, we follow Fryer (2019) in using a cluster-robust covariance estimator, clustering on precinct, and 5,000 samples were drawn for each force threshold and model specification.

A.6 Point identification of ATE

First, we note that strata sizes are identified with information on the total count of encounters by race (both reported and unreported).

$$
\begin{aligned}
\Pr(M_i(1) = 1, M_i(0) = 1, X_i = x) &= \Pr(M_i(1) = 1, M_i(0) = 1 | D_i = 0, X_i = x) \\
&= \Pr(M_i = 1 | D_i = 0, X_i = x) \\
\Pr(M_i(1) = 1, M_i(0) = 0, X_i = x) &= \Pr(M_i(1) = 1, X_i = x) \\
&\quad - \Pr(M_i(1) = 1, M_i(0) = 1, X_i = x)\Pr(M_i = 1 | D_i = 1, X_i = x) \\
&\quad - \Pr(M_i = 1 | D_i = 0, X_i = x) \\
\Pr(M_i(1) = 0, M_i(0) = 1, X_i = x) &= 0 \\
\Pr(M_i(1) = 0, M_i(0) = 0, X_i = x) &= 1 - \Pr(M_i(1) = 0, M_i(0) = 1, X_i = x) \\
&\quad - \Pr(M_i(1) = 1, M_i(0) = 0, X_i = x) \\
&\quad - \Pr(M_i(1) = 1, M_i(0) = 1, X_i = x) \\
&= 1 - \Pr(M_i = 1 | D_i = 1, X_i = x)
\end{aligned}
$$

We then reexpress the ATE in terms of strata-specific mean potential outcomes and simplify.

$$\text{ATE} = \mathbb{E}[Y_i(1, M_i(1))|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x]$$

$$\Pr(M_i(1) = 1, M_i(0) = 1|D_i = 1, X_i = x)\Pr(D_i = 1, X_i = x)$$

$$+ \mathbb{E}[Y_i(1, M_i(1))|D_i = 1, M_i(1) = 1, M_i(0) = 0, X_i = x]$$

$$\Pr(M_i(1) = 1, M_i(0) = 0|D_i = 1, X_i = x)\Pr(D_i = 1, X_i = x)$$

$$+ \mathbb{E}[Y_i(1, M_i(1))|D_i = 1, M_i(1) = 0, M_i(0) = 1, X_i = x]$$

$$\Pr(M_i(1) = 0, M_i(0) = 1|D_i = 1, X_i = x)\Pr(D_i = 1, X_i = x)$$

$$+ \mathbb{E}[Y_i(1, M_i(1))|D_i = 1, M_i(1) = 0, M_i(0) = 0, X_i = x]$$

$$\Pr(M_i(1) = 0, M_i(0) = 0|D_i = 1, X_i = x)\Pr(D_i = 1, X_i = x)$$

$$+ \mathbb{E}[Y_i(1, M_i(1))|D_i = 0, M_i(1) = 1, M_i(0) = 1, X_i = x]$$

$$\Pr(M_i(1) = 1, M_i(0) = 1|D_i = 0, X_i = x)\Pr(D_i = 0, X_i = x)$$

$$+ \mathbb{E}[Y_i(1, M_i(1))|D_i = 0, M_i(1) = 1, M_i(0) = 0, X_i = x]$$

$$\Pr(M_i(1) = 1, M_i(0) = 0|D_i = 0, X_i = x)\Pr(D_i = 0, X_i = x)$$

$$+ \mathbb{E}[Y_i(1, M_i(1))|D_i = 0, M_i(1) = 0, M_i(0) = 1, X_i = x]$$

$$\Pr(M_i(1) = 0, M_i(0) = 1|D_i = 0, X_i = x)\Pr(D_i = 0, X_i = x)$$

$$+ \mathbb{E}[Y_i(1, M_i(1))|D_i = 0, M_i(1) = 0, M_i(0) = 0, X_i = x]$$

$$\Pr(M_i(1) = 0, M_i(0) = 0|D_i = 0, X_i = x)\Pr(D_i = 0, X_i = x)$$

$$- \mathbb{E}[Y_i(0, M_i(0))|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x]$$

$$\Pr(M_i(1) = 1, M_i(0) = 1|D_i = 1, X_i = x)\Pr(D_i = 1, X_i = x)$$

$$- \mathbb{E}[Y_i(0, M_i(0))|D_i = 1, M_i(1) = 1, M_i(0) = 0, X_i = x]$$

$$\Pr(M_i(1) = 1, M_i(0) = 0|D_i = 1, X_i = x)\Pr(D_i = 1, X_i = x)$$

$$- \mathbb{E}[Y_i(0, M_i(0))|D_i = 1, M_i(1) = 0, M_i(0) = 1, X_i = x]$$

$$\Pr(M_i(1) = 0, M_i(0) = 1|D_i = 1, X_i = x)\Pr(D_i = 1, X_i = x)$$

$$- \mathbb{E}[Y_i(0, M_i(0))|D_i = 1, M_i(1) = 0, M_i(0) = 0, X_i = x]$$

$$\Pr(M_i(1) = 0, M_i(0) = 0|D_i = 1, X_i = x)\Pr(D_i = 1, X_i = x)$$

$$- \mathbb{E}[Y_i(0, M_i(0))|D_i = 0, M_i(1) = 1, M_i(0) = 1, X_i = x]$$

$$\Pr(M_i(1) = 1, M_i(0) = 1|D_i = 0, X_i = x)\Pr(D_i = 0, X_i = x)$$

$$- \mathbb{E}[Y_i(0, M_i(0))|D_i = 0, M_i(1) = 1, M_i(0) = 0, X_i = x]$$

$$\Pr(M_i(1) = 1, M_i(0) = 0|D_i = 0, X_i = x)\Pr(D_i = 0, X_i = x)$$

$$- \mathbb{E}[Y_i(0, M_i(0))|D_i = 0, M_i(1) = 0, M_i(0) = 1, X_i = x]$$

$$\Pr(M_i(1) = 0, M_i(0) = 1|D_i = 0, X_i = x)\Pr(D_i = 0, X_i = x)$$

$$- \mathbb{E}[Y_i(0, M_i(0))|D_i = 0, M_i(1) = 0, M_i(0) = 0, X_i = x]$$

$$\Pr(M_i(1) = 0, M_i(0) = 0|D_i = 0, X_i = x) \Pr(D_i = 0, X_i = x)$$

under mandatory reporting

$$= \mathbb{E}[Y_i(1, M_i(1))|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x]$$
$$\Pr(M_i(1) = 1, M_i(0) = 1|D_i = 1, X_i = x) \Pr(D_i = 1, X_i = x)$$
$$+ \mathbb{E}[Y_i(1, M_i(1))|D_i = 1, M_i(1) = 1, M_i(0) = 0, X_i = x]$$
$$\Pr(M_i(1) = 1, M_i(0) = 0|D_i = 1, X_i = x) \Pr(D_i = 1, X_i = x)$$
$$+ \mathbb{E}[Y_i(1, M_i(1))|D_i = 0, M_i(1) = 1, M_i(0) = 1, X_i = x]$$
$$\Pr(M_i(1) = 1, M_i(0) = 1|D_i = 0, X_i = x) \Pr(D_i = 0, X_i = x)$$
$$+ \mathbb{E}[Y_i(1, M_i(1))|D_i = 0, M_i(1) = 1, M_i(0) = 0, X_i = x]$$
$$\Pr(M_i(1) = 1, M_i(0) = 0|D_i = 0, X_i = x) \Pr(D_i = 0, X_i = x)$$
$$- \mathbb{E}[Y_i(0, M_i(0))|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x]$$
$$\Pr(M_i(1) = 1, M_i(0) = 1|D_i = 1, X_i = x) \Pr(D_i = 1, X_i = x)$$
$$- \mathbb{E}[Y_i(0, M_i(0))|D_i = 1, M_i(1) = 0, M_i(0) = 1, X_i = x]$$
$$\Pr(M_i(1) = 0, M_i(0) = 1|D_i = 1, X_i = x) \Pr(D_i = 1, X_i = x)$$
$$- \mathbb{E}[Y_i(0, M_i(0))|D_i = 0, M_i(1) = 1, M_i(0) = 1, X_i = x]$$
$$\Pr(M_i(1) = 1, M_i(0) = 1|D_i = 0, X_i = x) \Pr(D_i = 0, X_i = x)$$
$$- \mathbb{E}[Y_i(0, M_i(0))|D_i = 0, M_i(1) = 0, M_i(0) = 1, X_i = x]$$
$$\Pr(M_i(1) = 0, M_i(0) = 1|D_i = 0, X_i = x) \Pr(D_i = 0, X_i = x)$$

under mediator monotonicity

$$= \mathbb{E}[Y_i(1, M_i(1))|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x]$$
$$\Pr(M_i(1) = 1, M_i(0) = 1|D_i = 1, X_i = x) \Pr(D_i = 1, X_i = x)$$
$$+ \mathbb{E}[Y_i(1, M_i(1))|D_i = 1, M_i(1) = 1, M_i(0) = 0, X_i = x]$$
$$\Pr(M_i(1) = 1, M_i(0) = 0|D_i = 1, X_i = x) \Pr(D_i = 1, X_i = x)$$
$$+ \mathbb{E}[Y_i(1, M_i(1))|D_i = 0, M_i(1) = 1, M_i(0) = 1, X_i = x]$$
$$\Pr(M_i(1) = 1, M_i(0) = 1|D_i = 0, X_i = x) \Pr(D_i = 0, X_i = x)$$
$$+ \mathbb{E}[Y_i(1, M_i(1))|D_i = 0, M_i(1) = 1, M_i(0) = 0, X_i = x]$$
$$\Pr(M_i(1) = 1, M_i(0) = 0|D_i = 0, X_i = x) \Pr(D_i = 0, X_i = x)$$
$$- \mathbb{E}[Y_i(0, M_i(0))|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x]$$
$$\Pr(M_i(1) = 1, M_i(0) = 1|D_i = 1, X_i = x) \Pr(D_i = 1, X_i = x)$$
$$- \mathbb{E}[Y_i(0, M_i(0))|D_i = 0, M_i(1) = 1, M_i(0) = 1, X_i = x]$$

$$\Pr(M_i(1) = 1, M_i(0) = 1|D_i = 0, X_i = x) \Pr(D_i = 0, X_i = x)$$

under treatment ignorability

$$
\begin{aligned}
&= \mathbb{E}[Y_i(1, M_i(1))|D_i = 1, M_i(1) = 1, M_i(0) = 1, X_i = x] \Pr(M_i(1) = 1, M_i(0) = 1, X_i = x) \\
&\quad + \mathbb{E}[Y_i(1, M_i(1))|D_i = 1, M_i(1) = 1, M_i(0) = 0, X_i = x] \Pr(M_i(1) = 1, M_i(0) = 0, X_i = x) \\
&\quad - \mathbb{E}[Y_i(0, M_i(0))|D_i = 0, M_i(1) = 1, M_i(0) = 1, X_i = x] \Pr(M_i(1) = 1, M_i(0) = 1, X_i = x)
\end{aligned}
$$

which can be recovered from observed data

$$
\begin{aligned}
&= \mathbb{E}[Y_i|D_i = 1, M_i(D_i) = 1, X_i = x] \Pr(M_i = 1|D_i = 1, X_i = x) \\
&\quad - \mathbb{E}[Y_i|D_i = 0, M_i(D_i, X_i = x) = 1, X_i = x] \Pr(M_i = 1|D_i = 0, X_i = x)
\end{aligned}
$$

which reduces to Proposition 2 in the no-covariate case.

A.7 Derivation of outcome test bounds on ρ

Our paper focuses on the difficulty of estimating a race effect on post-stop police behavior such as the use of force. However, another popular approach, the outcome test, focuses on establishing whether there exists any bias in the decision to stop a civilian (Becker, 1971; Goel, Rao and Shroff, 2016; Engel, 2008; Knowles, Perisco and Todd, 2001; Ridgeway and MacDonald, 2010). Because the degree of the statistical bias we explore is a function of racial discrimination in stopping decisions, it is useful to clarify the assumptions undergirding outcome tests. In the process, we demonstrate that the principal stratification framework sheds light on the precise interpretation of outcome tests, and we prove that the outcome test can be used to establish a lower bound on the share of police stops of racial minorities that are racially discriminatory.

Outcome tests compare the rates of finding evidence of a crime—conditional on a suspect being stopped by police—across racial groups. The logic behind the test is that if the decision to stop a civilian is unbiased, the rate of discovering evidence of a crime ("hit rates") should be identical across groups. Proponents of outcome tests thus claim that differences in hit rates amount to evidence of racially biased policing. The empirical observation that hit rates are lower among minority stops can be written as $\mathbb{E}[Y_i|D_i = 0, M_i = 1] > \mathbb{E}[Y_i|D_i = 1, M_i = 1]$, where Y_i is an indicator, say, for finding contraband on a suspect. However, interpreting the above inequality as evidence of racial discrimination in fact requires assumptions that closely mirror those we describe above.

To see this, first observe that the overall hit rate among minority stops can be decomposed into

the weighted average of the hit rate among always-stop encounters and the hit rate among the (possibly nonexistent) set of racially discriminatory stops. In contrast, if we invoke Assumption 2 (which states that there are no white civilians stopped in circumstances where a minority civilian would be allowed to pass), then stops involving white civilians belong exclusively to the always-stop group.[1] In this case, the empirical difference in hit rates can be rewritten in the potential outcomes framework as

$$\mathbb{E}[Y_i(0,1)|M_i(1) = 1, M_i(0) = 1, D_i = 0]$$
$$-\mathbb{E}[Y_i(1,1)|M_i(1) = 1, M_i(0) = 1, D_i = 1]\Pr(M_i(1) = 1, M_i(0) = 1|D_i = 1, M_i = 1)$$
$$- \mathbb{E}[Y_i(1,1)|M_i(1) = 1, M_i(0) = 0, D_i = 1]\Pr(M_i(1) = 1, M_i(0) = 0|D_i = 1, M_i = 1) > 0 \quad (6)$$

A major critique of the outcome test is that observed racial disparities in hit rates alone do not constitute evidence of racially discriminatory stops because of the problem of "infra-marginality" (Ayres, 2002; Simoiu, Corbett-Davies and Goel, 2017). This critique suggests that the above inequality may hold simply because white civilians in always-stop encounters engage in more criminal conduct than minority suspects. In other words, the analyst might observe $\mathbb{E}[Y_i|D_i = 1, M_i = 1] < \mathbb{E}[Y_i|D_i = 0, M_i = 1]$ even if $\Pr(M_i(1) = 1, M_i(0) = 0) = 0$—that is, with no discrimination in stops—as long as $\mathbb{E}[Y_i(1,1)|M_i(1) = 1, M_i(0) = 1, D_i = 1] < \mathbb{E}[Y_i(0,1)|M_i(1) = 1, M_i(0) = 1, D_i = 0]$. Some analysts employing the outcome test cast this scenario as unlikely, arguing that absent racial bias in stopping, "it would be difficult to explain why…whites for some reason had a systematically higher chance of possessing evidence of illegality" (Ayres, 2002) (137) and "there are not compelling reasons to suspect" this to be the case (138). Indeed, the validity of the outcome test hinges on the assumption that white and minority civilians in always-stop encounters commit crimes at the same rates, or that $\mathbb{E}[Y_i(1,1)|M_i(1) = 1, M_i(0) = 1, D_i = 1] = \mathbb{E}[Y_i(0,1)|M_i(1) = 1, M_i(0) = 1, D_i = 0]$. This assumption closely parallels Assumption 4, which requires that treatment status is ignorable with respect to potential outcomes. (For simplicity, we suppose that this holds without conditioning on covariates, but the result also holds within levels of $X_i = x$.) In this case, the observed racial difference in hit rates can be rewritten as

$$\mathbb{E}[Y_i|D_i = 0, M_i = 1] - \mathbb{E}[Y_i|D_i = 1, M_i = 1]$$
$$= \Big(\mathbb{E}[Y_i(0,1)|M_i(1) = 1, M_i(0) = 1, D_i = 0]$$
$$-\mathbb{E}[Y_i(1,1)|M_i(1) = 1, M_i(0) = 0, D_i = 1]\Big)\Pr(M_i(1) = 1, M_i(0) = 0|D_i = 1, M_i = 1). \quad (7)$$

This formulation makes clear that the observed evidence gap is due to the difference in hit

---

[1]If we do not assume mediator monotonicity, and allow for the presence of stops of white suspects that would not have occurred if the suspect was a racial minority, then the inequality used to estimate the outcome test becomes uninformative with respect to racial discrimination.

rates between always-stop minority encounters—in which officers would also have stopped a white civilian—and racially discriminatory minority stops. If the former is assumed to produce more evidence of criminal behavior (Assumption 3; this might hold if racially discriminatory stops are made under weaker standards of evidence), then it can be seen from Equation 7 that the empirical difference in hit rates implies that $\Pr(M_i(1) = 1, M_i(0) = 0) > 0$: that there must exist encounters in which minority civilians would be stopped but white civilians would not, precisely as proponents of the outcome test suggest.

Equation 7 also shows that outcome tests are unable to identify the exact prevalence of racial stops. Outcome tests allow the analyst to infer whether there is *any* racial bias in the decision to stop a suspect—but only if the analyst makes assumptions similar to those we outline above. However, we show that the outcome test can *partially* identify a range of possible proportions of racial stops. This clarification is useful, as it allows us later in this analysis to appeal to a published study of hit rates (Goel, Rao and Shroff, 2016) to help characterize the statistical bias in analyses of post-stop police behavior (e.g. Fryer, 2019).

By rearranging Equation 7 and substituting observed quantities, we arrive at

$$\Pr(M_i(1) = 1, M_i(0) = 0|D_i = 1, M_i = 1) \;=\; \frac{\mathbb{E}[Y_i|D_i = 0, M_i = 1] - \mathbb{E}[Y_i|D_i = 1, M_i = 1]}{\mathbb{E}[Y_i|D_i = 0, M_i = 1] - \mathbb{E}[Y_i(1, 1)|M_i(1) = 1, M_i(0) = 0, D_i = 1]}$$

Although the second term in the denominator is unknown, the implied proportion of racially discriminatory stops is smallest when this value is zero—if, hypothetically, searches of racially stopped minorities never produce evidence. Thus, the outcome test suggests that *at least* ($\mathbb{E}[Y_i|D_i = 0, M_i = 1] - \mathbb{E}[Y_i|D_i = 1, M_i = 1]$) / $\mathbb{E}[Y_i|D_i = 0, M_i = 1]$ of all minority stops are racially discriminatory, and to the extent that racially discriminatory searches result in any evidence of contraband, the proportion could potentially be much larger.

B   Additional results

B.1   Coding schemes for dependent variables

In this section, we reanalyze the NYPD SQF data using both the original and revised coding schemes for dependent variables in Fryer (2019). In an analysis of the use of force by level of severity, Fryer (2019) codes binary outcomes indicating whether force is used at or above some threshold. However, rather than coding all encounters with lower levels of force than a given threshold as a zero, the analysis coded only encounters with no force at all as a zero, while levels of force between no force and the threshold level were dropped from the data.[2] This data dropping strategy, a form of selection on the dependent variable, is problematic. If the analyst suspects that civilian race affects which level of force is applied—the motivating hypothesis for this very analysis—then dropping data based on which level of force was applied is another form of post-treatment conditioning and will induce bias. Further, the amount of data lost under this coding scheme is substantial. In the case of the point-weapon threshold, for example, over one million encounters—over 20% of the data—appear to have been discarded despite containing sub-threshold force use, such as pushing a civilian to the ground. Table B1 displays the number of observations reported for various regressions in the original paper, our best attempts at replication, and the corrected procedure used in this paper.

We present results of our replication study using the original coding scheme and our corrected version side by side below. As the results show, a corrected analysis generally depresses the naïve treatment effects relative to the inadvisable coding scheme in Fryer (2019) and in most cases renders the original results statistically insignificant. However, these discrepancies in results across coding decisions do not alter the central point of our paper: post-treatment conditioning exerts a large downward bias on estimates of racially discriminatory uses of force.

---

[2]Fryer (2019) acknowledges this data dropping strategy, writing "To be clear, an observation that records only hands would be in the hands regression but <u>not</u> the regression which restricts the sample to observations in which individuals were at least forced to the ground," (21, emphasis in original).

Table B1: **Comparison of SQF Data Dimensions Based on Outcome Coding.** The table displays the number of observations from bivariate analyses of the use of force by the NYPD using three coding procedures for force outcomes. The first column displays the number of observations as reported in results in Fryer (2019) (Appendix Tables 3A-3G in original). The second column reports the number of observations we recover when using the coding procedure in Fryer (2019) which drops observations where some level of force was used that was below a given threshold. The third column displays the number of observations we recover when using our corrected coding procedure, which codes outcomes as a 1 if a certain force threshold is reached and 0 otherwise.

	N (published)	N (replicated)	N (corrected coding)
At least use of hands	4,927,962	4,980,701	4,980,701
At least push to wall	4,152,918	4,245,091	4,980,701
At least use of handcuffs	4,017,783	4,122,329	4,980,701
At least draw weapon	3,957,687	3,965,721	4,980,701
At least push to ground	3,950,324	3,958,374	4,980,701
At least point weapon	3,918,741	3,926,805	4,980,701
Use baton or pepper spray	3,900,977	3,909,064	4,980,701

Figure B1: **Replication of Fryer (2019) using various outcome coding rules.** The figure displays odds ratios generated by OLS regressions that show the effect of suspect race without covariates on the use of force across all force types generated using three approaches: the published OLS results from the Appendix of Fryer (2019) (black points and bars), our best attempt at replication of these results (blue points and bars), and results using our corrected outcome coding scheme (red points and bars). Revising the coding scheme so as to retain data on sub-threshold uses of force generally deflates estimated treatment effects.
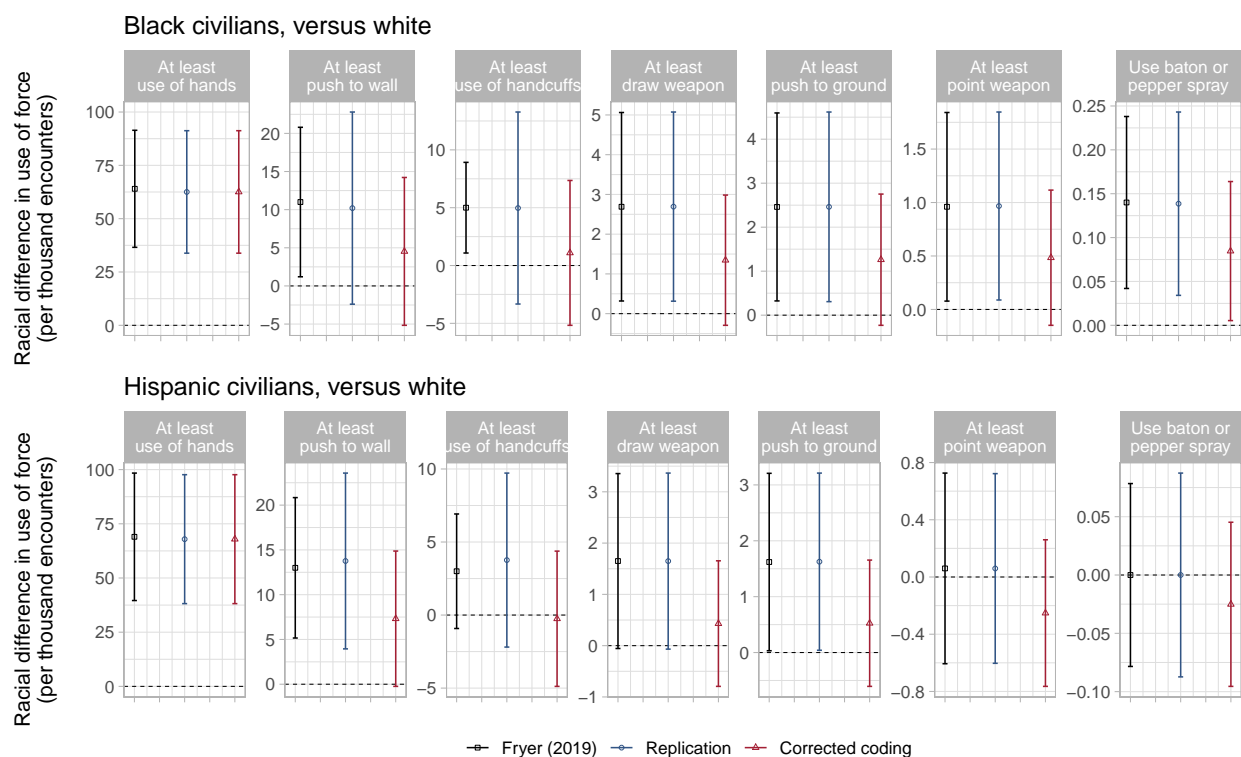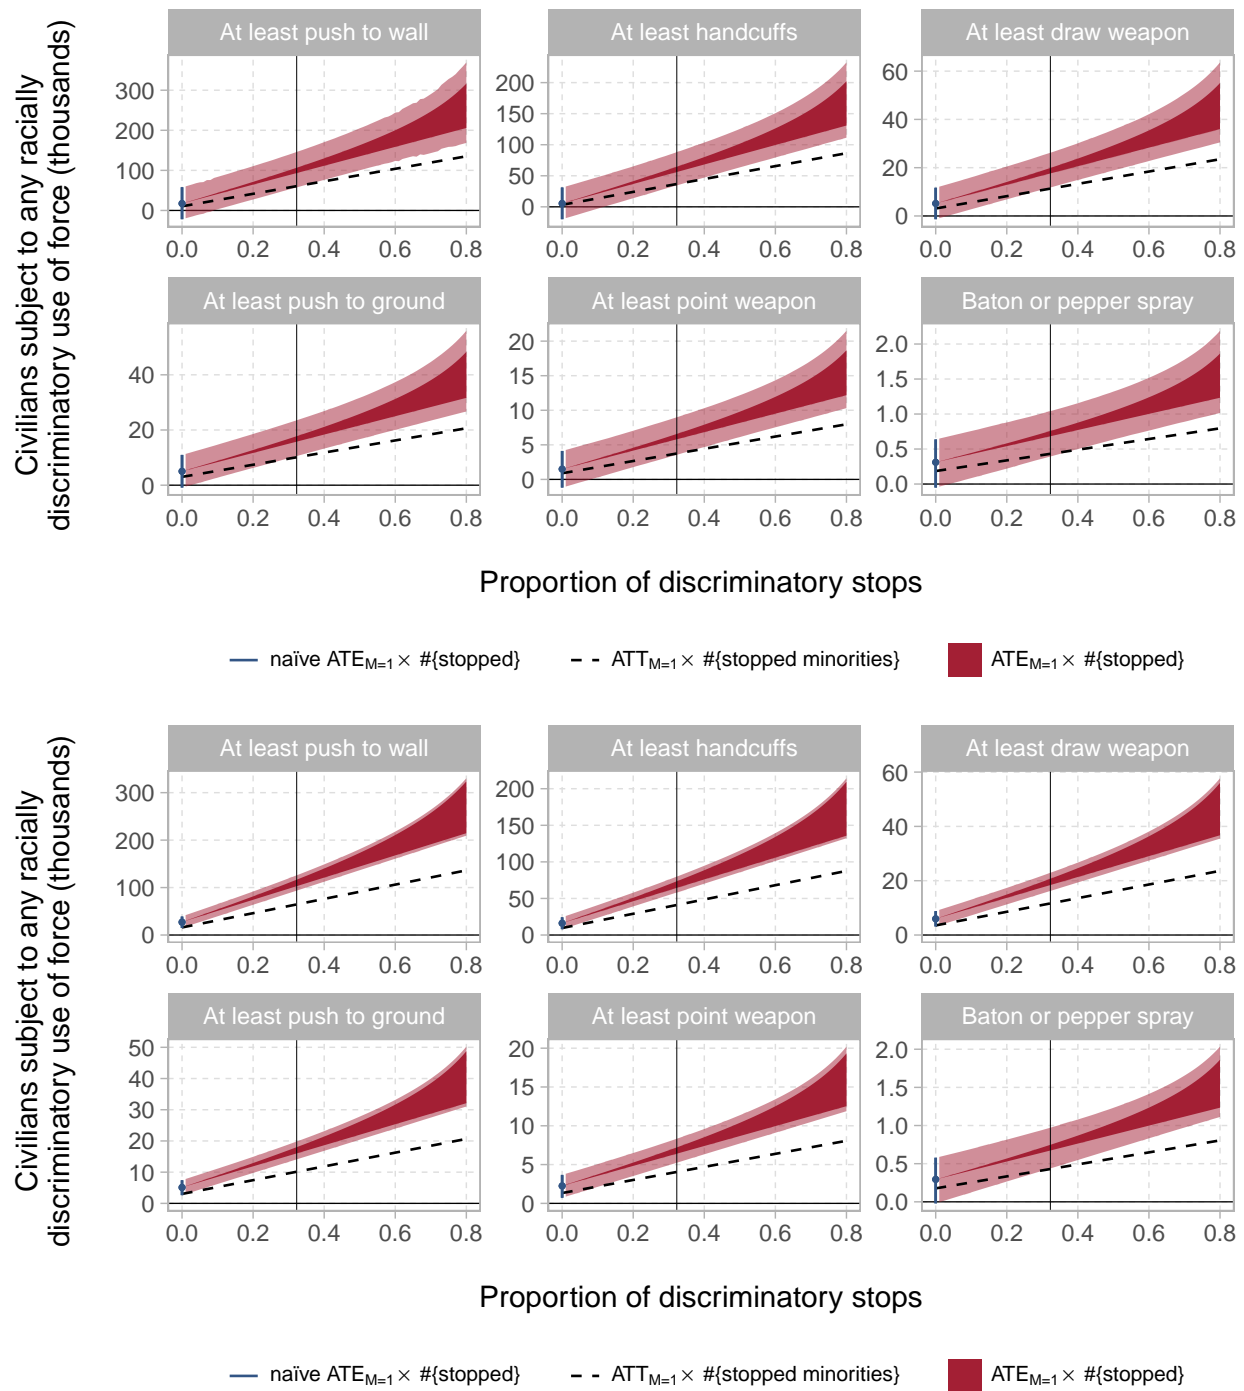
B.2 Varying levels of force

Figure B2: **Corrected ATE$_{M=1}$ and ATT$_{M=1}$ for encounters with Black and white civilians, varying levels of force.** This figure shows bounded effects comparing predicted levels of force when setting suspect race for all observations to black vs. white. These estimates use our corrected coding scheme for dependent variables (as described above). Results from regressions without covariates appear in the top panels and results from models with a full set of covariates appear in bottom panels.

Figure B3: **Corrected ATE$_{M=1}$ and ATT$_{M=1}$ for encounters with Hispanic and white civilians, varying levels of force.** This figure shows bounded effects comparing predicted levels of force when setting suspect race for all observations to Hispanic vs. white. These estimates use our corrected coding scheme for dependent variables (as described above). Results from regressions without covariates appear in the top panels and results from models with a full set of covariates appear in bottom panels.
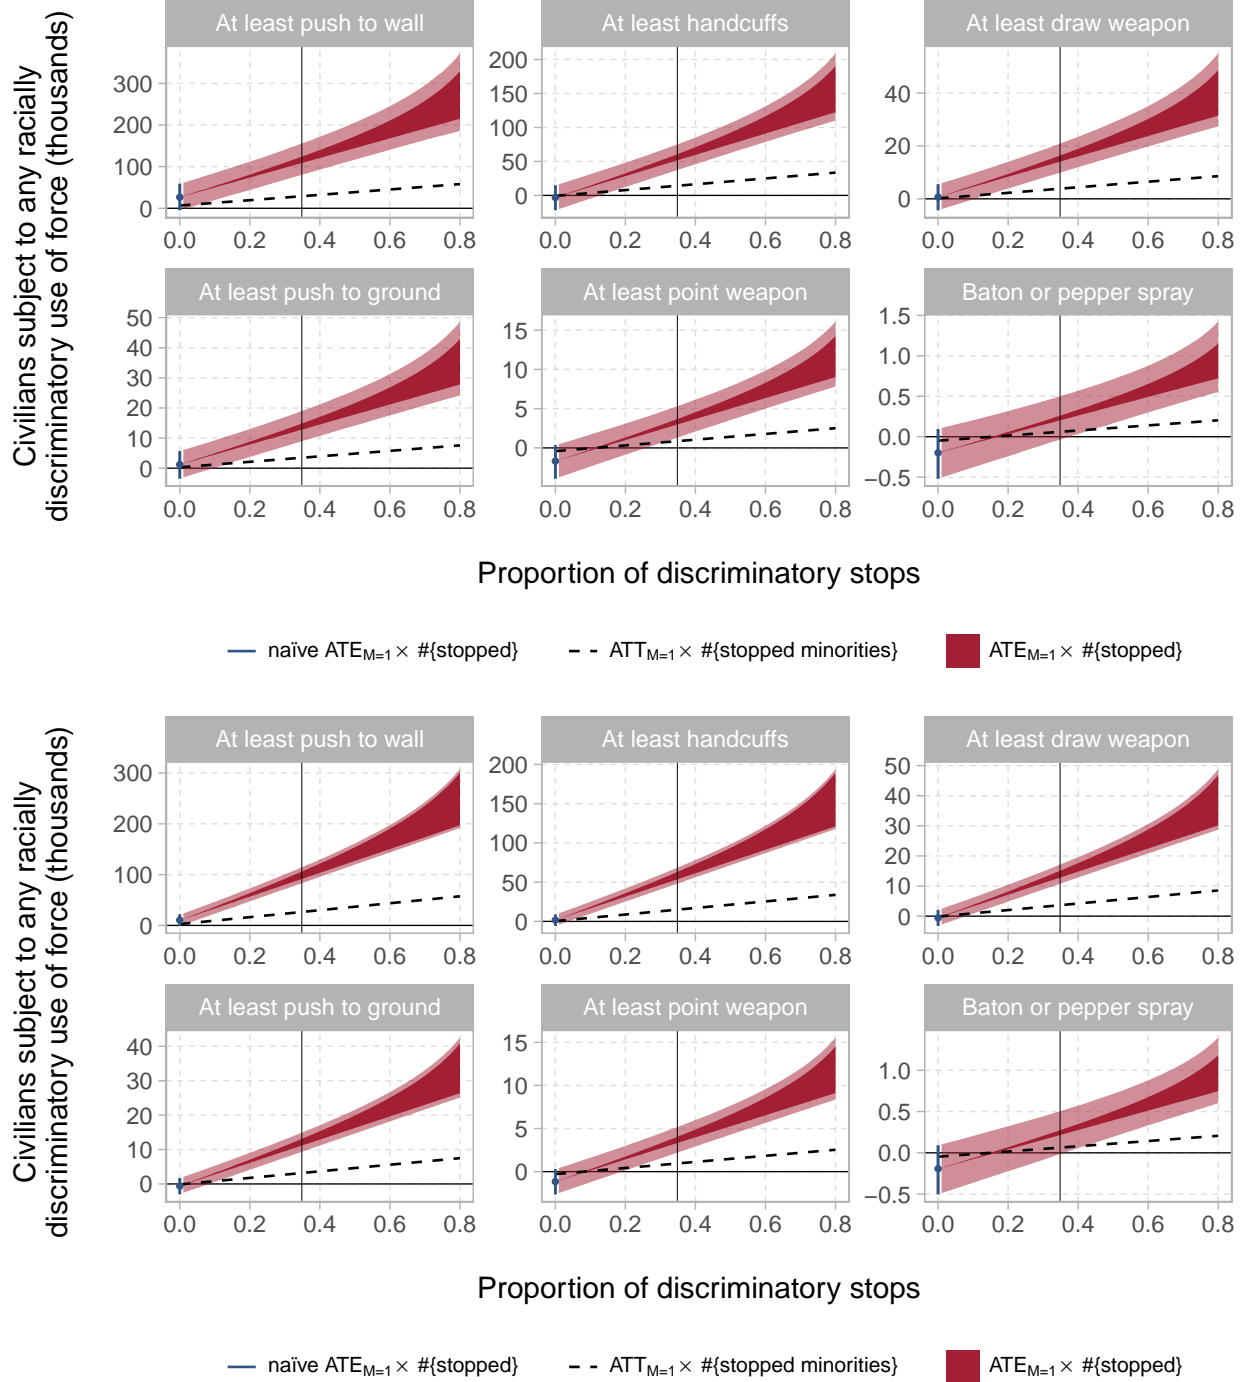
B.3 Excluding drug stops

Figure B4: **Bounds on race effect excluding drug stops.** This analysis replicates the analysis in Figure 4 in the main text excluding stops that were motivated by suspicion of a drug transaction, as such instances may violate the mediator monotonicity assumption. The results remain substantively similar.

B.4 Analysis of two races at a time

Figure B5: **Bounds on race effect limiting analysis to two racial groups of suspects.** Plots in the main text estimated bounds using data on multiple racial groups of suspects by predicting counterfactual values for every observation, regardless of a suspect's actual race, after model parameters were estimated. These figures reproduce the same analysis using only data on the two racial groups being compared, and exclude data on suspects who were not black, Hispanic or white entirely.

References

Ayres, Ian. 2002. "Outcome Tests of Racial Disparities in Police Practices." *Justice Research and Policy* 4(1-2):131–142.

Becker, Gary. 1971. *The Economics of Discrimination.* University of Chicago Press.

Engel, Robin. 2008. "A Critique of the "Outcome Test" in Racial Profiling Research." *Justice Quarterly* 25(1):1–36.

Fryer, Roland G. 2019. "An Empirical Analysis of Racial Differences in Police Use of Force." *Journal of Political Economy* 127(3):1210–1261.

Goel, Sharad, Justin M. Rao and Ravi Shroff. 2016. "Precinct or Prejudice? Understanding Racial Disparities in New York City's Stop-And-Frisk Policy." *Annals of Applied Statistics* 10(1):365–394.

Horowitz, Joel L. and Charles F. Manski. 2000. "Nonparametric Analysis of Randomized Experiments With Missing Covariate and Outcome Data." *Journal of the Americal Statistical Association* 95(449):77–84.

Knowles, J., N. Perisco and P. Todd. 2001. "Racial bias in motor vehicle searches: Theory and evidence." *Journal of Political Economy* 109(1):203–229.

Knox, Dean, Teppei Yamamoto, Matthew A. Baum and Adam J. Berinsky. 2019. "Design, Identification, and Sensitivity Analysis for Patient Preference Trials." *Journal of the American Statistical Association* 00(0):1–15.

Ridgeway, Greg and John MacDonald. 2010. *Race, Ethnicity, and Policing: New and Essential Readings.* NYU Press chapter Methods for Assessing Racially Biased Policing.

Robins, J.M. and S. Greenland. 1992. "Identifiability and exchangeability for direct and indirect effects." *Epidemiology* 3(2):143–155.

Simoiu, Camelia, Sam Corbett-Davies and Sharad Goel. 2017. "The problem of infra-marginality in outcome tests for discrimination." *The Annals of Applied Statistics* 11(3):1193–1216. https://arxiv.org/abs/1706.05678.